



NVIO

ABSTRACT

Details of the NVIO superscalar processor core including instruction set description.

Robert Finch

NVIO

©2019 Robert Finch

“If you build it, he will come” – Field of Dreams

Introduction

Features

- 80-bit data path, double-extended floating-point
- 64 entry general purpose register file
- 64 entry floating-point register file
- 3-way out-of-order (ooo) superscalar execution
- 40-bit instructions, three per 128-bit bundle
- Instruction L1, L2 and data L1, L2 caches
- 7 entry write buffer
- Dual memory channels

History

NVIO is a work-in-progress beginning in May 2019. NVIO originated from FT64 which originated from RiSC-16 by Dr. Bruce Jacob. RiSC-16 evolved from the Little Computer (LC-896) developed by Peter Chen at the University of Michigan. See the comment in FT64.v. NVIO is the author's second attempt at an 80 bit ISA. Other attempts including DSD9. The author has tried to be innovative with this design borrowing ideas from a number of other processing cores.

Design Rationale

It may seem strange to design an 80-bit cpu. Some of the rationale for doing so is double-extended floating-point is 80-bits. In order to support the double extended fp internal registers and busses need to be at least 80 bits in size. The author decided to limit things to 80 bits rather than go with a 128-bit design which would have been much larger. Given that a significant portion of the core is 80-bit it was decided just to make the whole core an 80-bit core. Double-extended floating-point offers good precision with which to do business apps. It's approximately 19 digits of precision which amounts to 13 digits plus 6 decimal digits for instance. Ordinary double-precision arithmetic only offers about 16 digits, which isn't quite enough for some business apps once things like rounding are considered. That'd only be 10+6 decimal points.

Nomenclature

The ISA refers to primitive object sizes following the convention suggested by Knuth of using Greek.

Number of Bits		Instructions
8	byte	LDB, STB
16	wyde	LDW, STW
32	tetra	LDT, STT
40	penta	LDP, STP
64	octa	LDO, STO
80	deci	LDD, STD

The register used to address instructions is referred to as the instruction pointer or IP register. The instruction pointer is a synonym for instruction pointer or PC register.

Development Aspects

Device Target

The core has been developed with FPGA usage in mind.

Implementation Language

The core is implemented in the System Verilog language primarily for its ability to process array objects. Much of the core is plain vanilla Verilog code.

Programming Model

Registers

Overview

The ISA is a 64-register machine with separate integer and floating-point register files. There are a large number of control and status (CSR) registers which hold an assortment of specific values relevant to processing. Why so many registers? With modern software growing ever more complex there are more statically assigned registers for performance reasons. There are about 20 registers with a static usage assignment. Were the design to include only 32 register that would leave only about 12 registers for general use, which may impact performance.

Due to the use of block memories to implement the register file, there is a bonus of sixteen different register sets, as the block memories have a minimum size associated with them. The design makes use of all the memory resources of the block memory. The availability of multiple sets of registers means that the register file may not need to be saved and restored during a thread switch.

Register Usage Convention


The register usage convention probably has more to do with software than hardware. Excepting a couple of special cases, the registers are general purpose in nature.

R0 always has the value zero. r61 is the link register used implicitly by the call instruction. The last two registers r62 and r63 are used for stack references and subject to stack bounds checking.

Register	Description / Suggested Usage	Saver
r0	always reads as zero	
r1-r4	return values / exception	caller
r5-r22	temporaries	caller
r23-r35	register variables	callee
r36-r43	function arguments	caller
r44-r50	reserved for system	
r51-r52	garbage collector	
r53-r54	assembler usage (constant building)	
r55	type number / function argument	caller
r56	class pointer / function argument	caller
r57	thread pointer	callee
r58	global pointer	
r59	exception SP offset	callee
r60	exception link register	callee
r61	return address / link register	callee
r62	base / frame pointer	callee
r63	stack pointer (hardware)	callee

Instruction Pointer

The instruction pointer identifies which instruction to execute. The instruction pointer increments as instructions are processed. The increment may be overridden using one of the flow control instructions. The instruction pointer addresses 40-bit instruction parcels located in a 128-bit bundle as a result the instruction pointer always has the last four bits one of 0h, 5h, or Ah. The instruction pointer register is also split into two sections. Branch instructions only update the lower 23-bits. Only the lower 40 bits of the IP increment.



Register Zero

Register zero – r0 – always reads as zero.

Stack and Frame Pointers

Although the stack and frame pointer registers may be used with any instruction the core has special hardware to detect stack bounds violations by either the stack pointer or frame pointer. The stack and frame pointer registers should be kept aligned on wyde boundaries. That is, they should be a multiple of two, which has the least significant bit as zero. There is currently no hardware in the core to enforce alignment.

Control and Status Registers

Overview

There are a large number of control and status registers. All bits of a CSR may not be implemented in hardware.

Control Register Zero (CSR #000)

This register contains miscellaneous control bits including a bit to enable protected mode.

Bit		Description
0	Pe	Protected Mode Enable: 1 = enabled, 0 = disabled
8 to 13		
16		
30	DCE	data cache enable: 1=enabled, 0 = disabled
32	BPE	branch predictor enable: 1=enabled, 0=disabled
34	WBM	write buffer merging enable: 1 = enabled, 0 = disabled
35	SPLE	speculative load enable (1 = enable, 0 = disable) (0 default)
36		
63	D	debug mode status. this bit is set during an interrupt routine if the processor was in debug mode when the interrupt occurred.

This register supports bit set / clear CSR instructions.

DCE

Disabling the data cache is useful for some codes with large data sets to prevent cache loading of values that are used infrequently. Disabling the data cache may reduce security risks for some kinds of attacks. The instruction cache may not be disabled. Enabling / disabling the data cache is also available via the cache instruction.

BPE

Disabling branch prediction will significantly affect the cores performance but may be useful for debugging. Disabling branch prediction causes all branches to be predicted as not-taken. No entries will be updated in the branch history table if the branch predictor is disabled.

WBM bit

Merging of values stored to memory may be disabled by setting this bit. On reset write buffer merging is disabled because it is likely desirable to setup I/O devices. Many I/O devices require updates to individual bytes by separate store instructions. (Write buffer merging is not currently implemented).

SPLE

Enabling speculative loads give the processor better performance at an increased security risk to meltdown attacks.

HARTID (0x001)

This register contains a number that is externally supplied on the hartid_i input bus to represent the hardware thread id or the core number. No core should have the value zero as the hartid.

TICK (0x002)

This register contains a tick count of the number of clock cycles that have passed since the last reset. Note that this register should not be used for precise timing as the processor's clock

frequency may vary for performance and power reasons. The TIME CSR may be used for wall-clock timing as it has its own timing source.

PCR Paging Control (CSR 0x003)

This register controls the paged memory management unit. A more detailed description is available under the section on memory management.

AEC Arithmetic Exception Control (CSR 0x004)

This register has controls to enable arithmetic exceptions and status bits to indicate the occurrence of exception conditions.

Exception Occurrence						Exception Enable					
63 37	36	35	34	33	32	31 5	4	3	2	1	0
	DIV	MUL	ASL	SUB	ADD		DIV	MUL	ASL	SUB	ADD

PMR Power Management Register (CSR 0x005)

This register contains bits to disable functional units in the core to conserve power. There is a bit for each functional unit.

			39 32	31 24	23 16	15 8	7 0
			FCU	MEM	FPU	ALU	Inst. Dec

It is not possible to disable all instruction decoders at the same time. There will always be at least one active decoder. Similarly, for the ALU; ALU #0 is always enabled. There will always be at least one memory channel on. The FCU can't be turned off, however the performance enhancing components can be. The branch predictor, branch target buffer and return stack buffer may all be turned off as a single unit. Power management features are not currently implemented.

CAUSE (0x006)

This register contains a code indicating the cause of an exception or interrupt. The break handler will examine this code in order to determine what to do. Only the low order 16 bits are implemented. The high order bits read as zero and are not updateable. The low order eight bits are loaded from the cause field of the BRK instruction. The next eight bits are loaded from the user₆ field of the break instruction.

BADADDR (CSR 0x007)

This register contains the effective address for a load / store operation that caused a memory management exception or a bus error. Note that the address of the instruction causing the exception is available in the EIP register.

PCR2 Paging Control (CSR 0x008)

This register controls the paged memory management unit. A more detailed description is available under the section on memory management.

Scratch (CSR 0x009)

This register is available for scratchpad use. It is typically swapped with a GPR during exception processing.

WBRC D (CSR 0x00A)

WBRC D stands for ‘write barrier record’. This register records the occurrence of a pointer store operation done by an instruction using a particular register set. There is a separate bit for each register set in the machine. This register is used by the garbage collector (GC) to determine if a scan should occur on a set of registers.

BAD_INSTR (CSR 0x00B)

This register contains a copy of the exceptioned instruction.

SEMA (CSR 0x00C) Semaphores

This register is available for system semaphore or flag use. The least significant bit is tied to the reservation address status input (rb_i). It will be set if a SWC instruction was successful. The least significant bit is also cleared automatically when an interrupt (BRK) or interrupt return (RTI) instruction is executed. Any one of the remaining bits may also be cleared by an RTI instruction. This could be a busy status bit for the interrupt routine. Bits in this CSR may be set or cleared with one of the CSRxx instructions. This register has individual bit set / clear capability.

Semaphore	Usage Convention
0	LDDR / STDC status bit
1	system garbage collection protector
2	system
3	input / output focus list
4	keyboard
5	system busy
6	memory management
7-79	currently unassigned

VM_SEMA (CSR 0x00D) Semaphores

This register is available for system semaphore or flag use.

KEYS – (CSR 0x00E)

This register contains the collection of keys associated with the process for the memory system. Each key is twenty bits in size. The register contains four keys. The combination of CSR 0x00E and 0x00F are searched in parallel for keys matching the one associated with the memory page.

79	60	59	40	39	20	19	0
key4		key3		key2		key1	

KEYS – (CSR 0x00F)

This register contains the collection of keys associated with the process for the memory system. Each key is twenty bits in size. The register contains four keys.

79	60	59	40	39	20	19	0
key8		key7		key6		key5	

TCB (CSR 0x010)

This CSR register is reserved for use as a pointer to a control block for the currently running thread.

WRS (CSR 0x011)

This register indicates which register sets are wired, or perpetually present in the core. Register sets beginning with register set #0 and progressing to the value supplied in this register are wired. There are only 16 register sets in the machine. Note that the general register file, and the floating-point register file each use a register set. Wiring too many register sets may make it difficult for the OS to switch threads.

FSTAT (CSR 0x014) Floating Point Status and Control Register

The floating-point status and control register may be read using the CSR instruction. Unlike other CSR's the control register has its own dedicated instructions for update. See the section on floating point instructions for more information.

Bit		Symbol	Description
79:44	~		reserved
43:40	PRC		Default precision
39:32	RGS		Register set
31:29	RM	rm	rounding mode
28	E5	inexe	- inexact exception enable
27	E4	dbzxe	- divide by zero exception enable
26	E3	underxe	- underflow exception enable
25	E2	overxe	- overflow exception enable
24	E1	invopxe	- invalid operation exception enable
23	NS	ns	- non standard floating point indicator
Result Status			
22		fractie	- the last instruction (arithmetic or conversion) rounded intermediate result (or caused a disabled overflow exception)
21	RA	rawayz	rounded away from zero (fraction incremented)
20	SC	C	denormalized, negative zero, or quiet NaN
19	SL	neg <	the result is negative (and not zero)
18	SG	pos >	the result is positive (and not zero)
17	SE	zero =	the result is zero (negative or positive)
16	SI	inf ?	the result is infinite or quiet NaN
Exception Occurrence			
15	X6	swt	{reserved} - set this bit using software to trigger an invalid operation
14	X5	inerx	- inexact result exception occurred (sticky)
13	X4	dbzx	- divide by zero exception occurred
12	X3	underx	- underflow exception occurred
11	X2	overx	- overflow exception occurred
10	X1	giopx	- global invalid operation exception – set if any invalid operation exception has occurred
9	GX	gx	- global exception indicator – set if any enabled exception has happened
8	SX	sumx	- summary exception – set if any exception could occur if it was enabled - can only be cleared by software
Exception Type Resolution			
7	X1T	cvt	- attempt to convert NaN or too large to integer
6	X1T	sqrtn	- square root of non-zero negative
5	X1T	NaNComp	- comparison of NaN not using unordered comparison instructions
4	X1T	infzero	- multiply infinity by zero
3	X1T	zerozero	- division of zero by zero
2	X1T	infdiv	- division of infinities
1	X1T	subinf	- subtraction of infinities
0	X1T	snanx	- signaling NaN

DBAD_x (CSR 0x018 to 0x01B) Debug Address Register

These registers contain addresses of instruction or data breakpoints.

79	0
Address 79..0	

DBCR (CSR 0x01C) Debug Control Register

This register contains bits controlling the circumstances under which a debug interrupt will occur.

bits			
3 to 0	Enables a specific debug address register to do address matching. If the corresponding bit in this register is set and the address (instruction or data) matches the address in the debug address register then a debug interrupt will be taken.		
17, 16	This pair of bits determine what should match the debug address register zero in order for a debug interrupt to occur.		
	17:16		
	00	match the instruction address	
	01	match a data store address	
	10	reserved	
	11	match a data load or store address	
19, 18	This pair of bits determine how many of the address bits need to match in order to be considered a match to the debug address register. These bits are ignored when matching instruction addresses, which are always half-word aligned.		
	19:18		Size
	00	all bits must match	byte
	01	all but the least significant bit should match	char
	10	all but the two LSB's should match	tetra
	11	all but the three LSB's should match	octa
23 to 20	Same as 16 to 19 except for debug address register one.		
27 to 24	Same as 16 to 19 except for debug address register two.		
31 to 28	Same as 16 to 19 except for debug address register three.		
55 to 62	These bits are a history stack for single stepping mode. An exception will automatically disable single stepping mode and record the single step mode state on stack. Returning from an exception pops the single step mode state from the stack.		
63	This bit enables SSM (single stepping mode)		


DBSR (CSR 0x01D) - Debug Status Register

This register contains bits indicating which addresses matched. These bits are set when an address match occurs and must be reset by software.

bit	
0	matched address register zero
1	matched address register one
2	matched address register two
3	matched address register three
79 to 4	not used, reserved

CAS (CSR 0x02C) Compare and Swap

This register is to support the compare and swap (CAS) instruction. If the value in the addressed memory location identified by the CAS instruction is equal to the value in the CAS register, then the source register is written to the memory location, and the source register is loaded with the value 1. Otherwise if the value in the addressed memory location doesn't match the value in this register, then value at the memory location is loaded into the CAS register, and the source register is set to zero. No write to memory occurs if the match fails.



TVEC (0x030 to 0x033)

These registers contain the address of the exception handling routine for a given operating level. TVEC[0] (0x030) is used directly by hardware to form an address of the interrupt routine. The lower eight bits of TVEC[0] are not used. The lower bits of the interrupt address are determined from the operating level. TVEC[1] to TVEC[3] are used by the REX instruction.

CMDPARM (0x038 to 0x03D)

These registers are used to pass command parameters to the operating system.

IM_STACK (0x040)

This register contains the interrupt mask stack. When an exception or interrupt occurs, this register is shifted to the left and the low order bits are set to all ones, when an RTI instruction is executed this register is shifted to the right. On RTI the last stack entry is set to fifteen masking all interrupts on stack underflow. The low order four bits represent the current interrupt mask level. Only the low order 40 bits of the register are implemented.

OL_STACK (0x041)

This register contains the operating and data level stack. When an exception or interrupt occurs, this register is shifted to the left, when an RTI instruction is executed this register is shifted to the right. On RTI the last stack entry is set to zero which will select the machine operating level on stack underflow. The low order 20 bits are the code/stack operating level, bits 32 to 51 are the data operating level. Note there is extra unused space in this register to accommodate a change to more threads or greater stack depth. Bit 0,1 represent the current operating level. Bits 32,33 represent the current data level.

PL_STACK (0x042)

This register contains the privilege level stack. When an exception or interrupt occurs, this register is shifted to the left, when an RTI instruction is executed this register is shifted to the right. On RTI the last stack entry will be set to zero which will select privilege level zero on stack underflow. The low order eight bits of the register represent the current privilege level.

RS_STACK (0x043)

This register contains the register set selection stack. When an exception or interrupt occurs, this register is shifted to the left and register set #0 is selected, when an RTI instruction is executed this register is shifted to the right. On RTI the last stack entry will be set to zero which will select register set #0 on stack underflow.

STATUS (0x044)

This register contains the interrupt mask, operating level, and privilege level.

Bitno	Field	Description
0 to 3	IM	active interrupt mask level
4 to 5	~	reserved
6 to 13	PL	privilege level
14 to 19	RS	register set selection – general purpose registers, this also controls which bounds register set is viewable in the CSRs.
20 to 21	~	reserved
24 to 27	Thrd	active thread
28 to 31	IRQ	The level of interrupt that caused the hardware BRK.
32	VCA	indicates that vector chaining was active prior to an exception
40 to 47	ASID	active address space identifier
48 to 49	FS	floating point state
50 to 51	XS	additional core extension state
55	MPRV	memory privilege: This bit when true (1) causes memory operations to use the first stack privilege level when evaluating privilege and protection rules. (Bits 0 to 13 in the status reg).
56 to 60	VM	These bits control virtual memory options. Note that multiple options may be present at the same time. At reset all the bits are set to zero.
63	SD	

VM₅

Bit	Indicates	
0	1 = single bound	
1	1 = separate program and data bounds	
2	1 = lot protection system	
3	1 = simplified paged unit	
4	1 = paging unit	

BRS_STACK (0x046)

This register contains the register set selection stack for base and bounds registers. When an exception or interrupt occurs, this register is shifted to the left and the current status copied to the low order bits, when an RTI instruction is executed this register is shifted to the right and the status bits copied from the low order bits of the register. On RTI the last stack entry will be set to eight which will select register set #8 on stack underflow.

Normally the brs_stack matches the rs_stack, but it's convenient to have a different register for interrupt processing so that base and bounds register may be modified without switching the general-purpose register set.

EIP (0x048 to 0x051)

This sets of registers contains the interrupt or exception stack of the instruction pointer register. The top of the stack is register 0x48. When an interrupt or exception occurs register 0x48 to 0x50 are copied to the next register and the instruction pointer is placed into register 0x48. When an RTI instruction is executed the instruction pointer is loaded from register 0x048 and registers

0x048 to 0x050 are loaded with the next register. Register 0x051 is loaded with the address of the break handler so that in the event of an underflow the break handler will be executed.

CODEBUF (0x080 to 0x0BF)

This register range is for access to 64 adaptable code buffers. The code buffers are used by the EXEC instruction to execute code which may change at run-time. Bits 0 to 39 of the code buffer are used to encode the instruction. Bits 40 to 46 of the code buffer are used to encode the template associated with the instruction. See the EXEC instruction for more details.

IQ_CTR (0xFC0)

This register contains a 40-bit count of the number of instructions queued since the last reset.

BM_CTR (0xFC1)

This register contains a 40-bit counter of the number of branch misses since the last reset.

IRQ_CTR (0xFC3)

This register is reserved to contain a 40-bit count of the number of interrupt requests.

BR_CTR (0xFC4)

This register contains a 40-bit counter of the number of branches committed since the last reset.

TIME (0xFE0)

The TIME register corresponds to the wall clock real time. This register can be used to compute the current time based on a known reference point. The register value will typically be a fixed number of seconds offset from the real wall clock time. The lower 40 bits of the register are driven by the tm_clk_i clock time base input which is independent of the cpu clock. The tm_clk_i input is a fixed frequency used for timing that cannot be less than 10MHz. The low order 40 bits represent the fraction of one second. The upper 40 bits represent seconds passed. For example, if the tm_clk_i frequency is 100MHz the low order 40 bits should count from 0 to 99,999,999 then cycle back to 0 again. When the low order 40 bits cycle back to 0 again, the upper 40 bits of the register is incremented. The upper 40 bits of the register represent the number of seconds passed since an arbitrary point in the past.

Note that this register has a fixed time basis, unlike the TICK register whose frequency may vary with the cpu clock. The cpu clock input may vary in frequency to allow for performance and power adjustments.

INSTRET (0xFE1)

This register contains a count of the number of instructions retired (successfully completed) by the core.

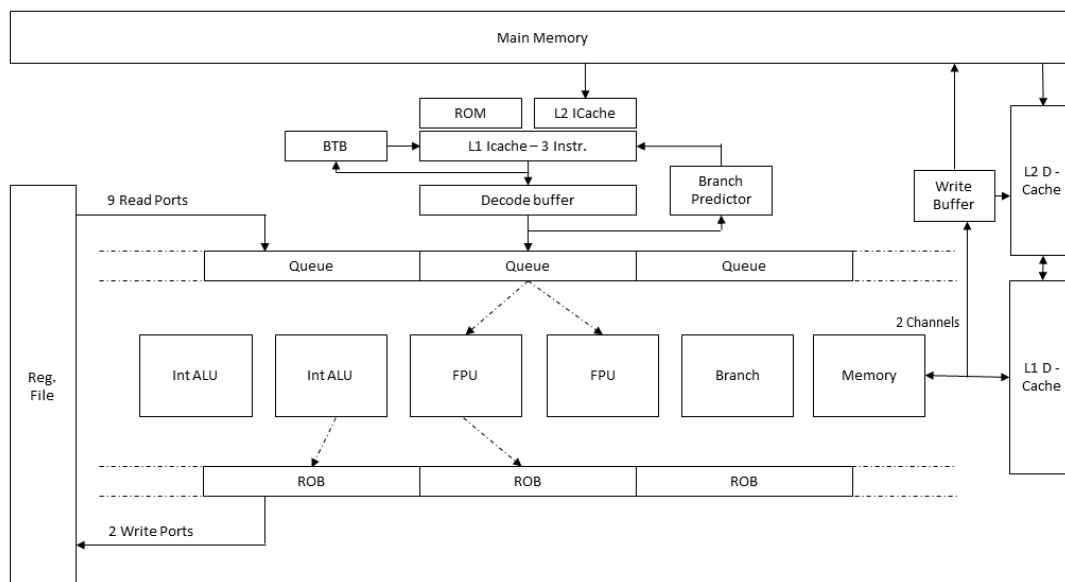
INFO (0xFF0 to 0xFFF)

This set of registers contains general information about the core including the manufacturer name, cpu class and name, and model number.

Pipeline

The pipeline is a typical superscalar pipeline as can be seen in the diagram below. A ROM with a two cycle latency feeds directly into the L1 cache. Up to three instructions are decoded and queued in a single cycle. After the instruction is queued it is assigned a re-order buffer entry. Once an instruction is finished results are transferred to the re-order buffer. Re-order buffer results are transferred to the register file in program order.

Pipeline Diagram



Caches

Overview


The core has both instruction and data caches to improve performance. Both caches are a two-level caches (L1, L2) allowing better performance. The first level cache is four-way set associative, the second level cache is also four-way set associative. The author initially had the first level cache fully associative based on a cam memory but found the resource requirements for the cam memory to be too large. The cache sizes of the instruction and data cache are available for reference from one of the INFO CSR registers.

Instructions

Since the instruction format affects the cache design it is mentioned here. For this design instructions are a single size and fit into 128-bit bundles. Specific formats are listed under the instruction set description section of this book.

L1 Instruction Cache

L1 is 2kB in size and made from distributed ram to get single cycle read performance. L1 is organized as 64 lines of 32-bytes. The following illustration shows the L1 cache organization for NVIO.



A 64-line cache was chosen as that matches the inherent size of single distributed ram component in the FPGA. It is the author's opinion that it would be better if the L1 cache were larger because it often misses due to its small size. However, using a larger distributed ram means going outside of the single lookup-table (LUT) and may then affect the clock cycle time. In short, the current design is an attempt to make it easy for the tools to create a fast implementation.

Note that supporting interrupts and cache misses, a requirement for a realistic processor design, adds complexity to the instruction stream. Reading the cache ram, selecting the correct instruction word and accounting for interrupts and cache misses must all be done in a single clock cycle.

While the L1 cache has single cycle reads it requires two clock cycles to update (write) the cache. The cache line to update needs to be provided by the tag memory which is unknown until after the tag updates.


One may have noticed the cache line size of 258 bits, the line size provides for two instruction bundles plus two bits which record error (tlb miss) state. One may also notice the lack of an instruction aligner, which is not required since instructions are fixed in size and processed as bundles.

L2 Instruction Cache

L2 is 16kB in size implemented with block ram. L2 is organized as 512 lines of 32 bytes. There is more flexibility in the design of the L2 cache since it's made up of block ram components. Once again, the cache is too small in the author's opinion, but it represents a trade-off in use of block ram resources for the cpu core versus using the block ram for other purposes such as memory management or data cache. There are only so many block rams in the FPGA.

The L2 cache has a read latency of three clock cycles to try and get the best clock cycle time out of the cache. It feeds the L1 cache with a cache-line wide bus so that only a single transfer cycle is required to update the L1 cache.

L2 Cache Organization



A cache load reads two bundles. Cache line size is 32 bytes.

An input register is used to feed the L1 cache directly on a load so that it doesn't have to wait for a read of the L2 cache before proceeding, otherwise there would be an additional three cycle delay during a cache miss.

If there is an error during the fetch, the L1 cache line is loaded with a break instruction corresponding to the error. A subsequent instruction fetch will cause the processor to execute the error routine.

The L2 cache is 16kB (512 rows * 32 bytes) composed of block ram which has a three cycle read latency. The L2 cache is four-way set associative.

The L2 cache reads two words from memory on a cache line load. While data is loading into the L2 cache an input register is also loaded with the data, the input register is transferred to the L1 cache, this is done so that the L1 cache update doesn't have to wait for the three cycle read latency of the L2 cache. Also fed into the input register are instructions for error processing should an error occur during the cache load.

Data Cache

The data cache organization is similar to the instruction cache with the exception that the data cache needs to be able to accept data written by store instructions as well as providing read data. The data cache also supports unaligned data access. These features make the data cache more complex than the instruction cache. Some of the latency of the data cache can be hidden by the presence of non-memory operating instructions in the instruction queue and the ability of the core to overlap data fetches.

The data cache has two read ports allowing two load operations to be in progress at the same time. The policy for stores is write-through. Stores always write through to memory. Since stores follow a write-through policy the latency of the store operation depends on the external memory system.

L1 Data Cache

Like the L1 instruction cache, the L1 data cache is 2kB in size, making use of 64 deep LUT rams. The data cache line size is 333 bits. 328 bits are required to support unaligned data access where the maximum overflow from a 256-bit region is 72 bits. Three bits are used to record error state. Like the I-cache the D-cache is four way set associative.

L2 Data Cache

The L2 data cache is organized as 512 lines of 32 bytes (16kB) and implemented with block ram. Access to the L2 data cache is multicycle. It isn't critical that the cache be able to update in single cycle as external memory access is bound to take many more cycles than a cache update. There is only a single write port on the data cache.

Cache Enables

The instruction cache is always enabled to keep hardware simpler and faster. Otherwise an additional multiplexor and control logic would be required in the instruction stream to read from external memory.

For some operations it may be desirable to disable the data cache so there is a data cache enable bit in control register #0. This bit may be set or cleared with one of the CSR instructions.

Cache Validation

A cache line is automatically marked as valid when loaded. The entire cache may be invalidated using the CACHE instruction. Individual cache lines may also be invalidated using the CACHE instruction.

Uncached Data Area

The address range \$F...FDxxxxx is an uncached 1MB data area. This area is reserved for I/O devices. The data cache may also be disabled in control register zero.

Write Buffering

The core has a configurable seven entry write buffer to enhance performance of store operations. The core always writes-through to memory and at the same time the cache is updated if a cache hit occurs. Loads will not occur until after the write buffer is emptied to ensure that data loaded isn't stale.

Load Bypassing

The write buffer does not currently support load bypassing of values contained in the write buffer to the load unit. Load bypassing has a higher security risk.

Branch Predictor

The branch predictor is a (2, 2) correlating predictor. The branch history is maintained in a 512-entry history table. It has three read ports for predicting branch outcomes, one port for each instruction in the fetch buffer. The branch predictor may be disabled by a bit in control register zero. When disabled all branches are predicted as not taken.

To conserve hardware the branch predictor uses a fifo that can queue up to two branch outcomes at the same time. Outcomes are removed from the fifo one at a time and used to update the branch history table which has only a single write port. In an earlier implementation of the branch predictor, two write ports were provided on the history table. This turned out to be relatively large compared to its usefulness.

Correctly predicting a branch turns the branch into a single-cycle operation. Branches are detected in the instruction fetch stage. During execution of the branch instruction the branch status is checked against the predicted status and if the two differ then a branch miss occurs. The miss address may be the branch's target address if the branch was supposed to be taken, otherwise it will be the address of the next instruction. If there was a branch miss, then the queue will be flushed, and new instructions loaded from the correct program path.

Branches

Target Address

Branches use absolute target addressing. The branch target field of the instruction is large enough that relative addressing is not required. Branching modifies the low order 23 bits of the instruction pointer while leaving the remaining bits as they are. Branching then takes place within an 8MB page of memory.

Performance

Branches should target the first slot (slot0) of a bundle for maximum performance. The core always fetches instructions at bundle aligned addresses. Branching into the middle of a bundle lowers performance because effectively fewer instructions are available for queueing. Branches should also be placed in the last slot (slot2) for maximum performance.

Clock Cycles

Under absolutely ideal conditions, 0.33 clocks are required to perform a branch. Ideal conditions being the branch is located in instruction slot #2 (so that the prior two slots may be executing instructions). Also register values must already be valid and the branch prediction is correct. Branches are predicted and taken in the fetch stage of the processor.

The predictor is about 90% accurate. So, with an eight-entry instruction queue on average about ½ the queue has to be flushed or 4 entries. That occurs 10% of the time. Also, assuming the branch is not in the last slot, about 1.5 slots are wasted on average. So, an average number of clock cycles for a branch would be about $(0.10 * 4 + 0.33) * 1.5 = 1.1$ clock cycles (or 1 if rounding).

Sequence Numbers

Sequence numbers are mentioned here because they are used by branch instructions to determine what to invalidate. There are a few different ways to handle invalidating instructions in the queue that follow a branch miss. Perhaps the lowest cost and fastest means is to use branch tags. Two other means are using a set of branch shadow bitmasks and using instruction sequence numbers.

Branch Tags

As branch instructions enter the queue, they are assigned a branch tag usually two bits. Non-branch instructions inherit the branch tag from the preceding branch. When a branch miss occurs all the instructions with the branch tag matching the missed branch are invalidated. When the branch commits to the state of the machine, the branch tag is freed up. Since there are only two bits used only four different tags are possible, meaning the processor can only speculate past four branches. In practical terms this isn't really a severe limitation.

Branch Shadow Masks

A means which the author came up. A shadow bit is allocated in the instruction queue for each branch. For instance, to support speculating past four branches four bits are required in the instruction queue. As instructions are entered into the queue the bit is set for the preceding branch instruction. When a branch miss occurs all the instructions with the shadow bit set for that branch are invalidated. The shadow bits are available to be reassigned to new branches when the branch commits to the machine state.

Sequence Numbers

Sequence numbers are probably conceptually the simplest means to understand. The order of instructions in the queue is associated with a number. The core assigns a sequence number to each instruction as it enters the instruction queue. One purpose of the sequence number is to allow the core to determine which instructions should be invalidated because of a branch miss. All the instructions in the queue with a sequence number coming after the branch instruction's sequence number will be invalidated. The sequence number assigned is the next highest number above that which is already in the queue. As instructions commit to the machine state, the sequence numbers of following instructions are decremented by the value of the sequence number of the committing instruction. This keeps the sequence numbers within range of the number of queue entries while maintaining the ordering relationship between the numbers.

The sequence number mechanism allows the core to speculate across any number of branches that might be in the instruction queue. However, sequence numbers require more bits in the instruction queue than branch tags and even more hardware. Since every instruction in the queue needs a number the number of bits required for the sequence number depends on the queue size. For a fifteen-entry queue a five-bit sequence number is required.

For debugging purposes, it can be handy to see the sequence number assigned to the instruction.

Branch Target Buffer (BTB)

The core has a 1k entry branch target buffer for predicting the target address of flow control instructions where the address is calculated and potentially unknown at time of fetch. Instructions covered by the BTB include jump-and-link, interrupt return and breakpoint instructions and branches to targets contained in a register.


Return Address Stack Predictor (RSB)

There is an address predictor for return addresses which can in some cases eliminate the flushing of the instruction queue when a return instruction is executed. The RET instruction is detected in the fetch stage of the core and a predicted return address used to fetch instructions following the return. JAL instructions using the link register as the source are also treated as return instructions. The return address stack predictor has a stack depth of 32 entries. On stack overflow or underflow, the prediction will be wrong, however performance will be no worse than not having a predictor. The return address stack predictor checks the address of the instruction queued following the RET against the address fetched for the RET instruction to make sure that the address corresponds.

Instruction Queue (ROB)

The instruction queue is a multi-entry re-ordering buffer (ROB). The size of the instruction queue is configurable between 4 and 16 entries. The instruction queue tracks an instructions progress and provides a holding place for operands and results. Each instruction in queue may be in one of a number of different states. The instruction queue is a circular buffer with head and tail pointers. Instructions are queued onto the tail and committed to the machine state at the head.

Instruction Queue – Re-order Buffer



The instruction queue is circular with eight slots. Each slot feeds a multiplexor which in turn feeds a functional unit. Providing arguments to the functional unit is done under the use of issue logic. Output from the functional unit is fed back to the same queue slot that issued to the functional unit. The queue slots are fed from the fetch buffers.

Queue Rate

Up to three instructions may queue during the same clock cycle depending on the availability of queue slots.

Sequence Numbers

The queue maintains a five-bit instruction sequence number which gives other operations in the core a clue as to the order of instructions. The sequence number is assigned when an instruction queues. The number assigned is one more than the maximum sequence number found in the queue. As instructions commit, the sequence numbers of all the remaining queue entries are decremented. This keeps the sequence number within five bits. Branch instructions need to know when the next instruction has queued in order to detect branch misses. The instruction pointer cannot be used to determine the instruction sequence because there may be a software loop at work which causes the instruction pointer to cycle backwards even though it's really the next instruction executing.

Queueing of Flow Control Operations

Flow control operations are not done until sometime after the next instruction queues. This is necessary to determine address miss-predicts during the flow control operation. Waiting until the next instruction queues avoids the problem of false mis-predictions. A consequence of waiting for the next instruction to queue is that flow control operations may only issue from one of the first

seven queue slots relative to the head of the queue. Note however that if the instruction queue is full the flow control operation will issue anyway otherwise the core could become deadlocked. When the core issues a flow control operation because the queue is full it will most likely cause a branch-miss state, which may reduce performance.

Issue Logic

Issue logic is responsible for assigning instructions to functional units. Instructions cannot be issued unless all operands are available, and the functional unit is also available.

The amount of issue logic required grows at a more than linear rate corresponding to the number of queue entries in the re-order buffer. This is in part due to the need for the issue logic to be synchronous in nature for an FPGA. There are more elegant ways to implement the issue logic using asynchronous loops, however these are not possible with an FPGA implementation. The amount of issue logic may be reduced by a core configuration define at some loss of performance. Since instructions that have just queued at the tail of the queue are unlikely to be ready to be processed the issue logic for those queue entries can be omitted. Instructions more towards the head of the queue will be more likely to be ready to issue.

Issue Rate

The functional units of NVIO include two alu's, two floating-point units, a memory unit, and a flow control unit (branch unit). Instructions may be issued to any and all functional units during a single clock cycle. The memory unit can handle two requests at the same time. As a result, up to seven instructions may be issued in a single clock cycle. In practice fewer instructions will be ready to be issued. The author has noted, with limited testing, that issuing a third memory operation in the same clock cycle is rarely done. The memory unit is typically 2/3 occupied or less.

Execute Logic

Instructions are executed on functional units after they have been issued to the unit. The execution logic for NVIO consists of two alu's, two address generators, two floating-point units, a flow control unit and a memory unit.

ALU's

Most instructions execute on one of two alu's. The alu's are asymmetrical. The first alu supports all operations including rarely performed operations, the second alu supports a subset of the operations which represents the most commonly performed operations. Splitting the functionality like this allows the core to support a wide variety of instructions while at the same time not using too many resources for rarely used operations. The issue logic knows about the difference in the ALU's and will issue what it can to the second alu if the first is busy. It's best to intermix commonly used instructions with rarely used ones to keep both alu's busy.

Shift Operations

The immediate constant for a shift is limited to six bits or the values 0 to 63. If it's required to shift by more than 63 bits two shift operations will be needed, or the value of the shift may come from a register.

Fixed Point Arithmetic

The core features the capability to perform fixed point arithmetic including addition, subtraction, multiplication and division. The point position is at $\frac{1}{2}$ the machine width or 40 binary point positions.

AGEN's


The core supports dual address generators which compute memory addresses. The address generators handle both register indirect with displacement and indexed register indirect addressing modes. The stack address for push operations and the effective address for the LEA instruction are also computed by the agen.

Floating Point Unit


The floating-point unit is used to execute almost all floating-point operations. The exception is floating point branches which are executed on the flow control unit. The length of time required to complete a floating-point operation varies depending on the instruction.

It would be better for performance if the floating-point unit were broken apart into several separate units. Each unit would have associated issue logic. That way the pipelining of the individual operations could be put to better use. Some machines have the fp multiplier separate from fp addition/subtraction and other units. This could be done but would require more resources from the FPGA. Having multiple floating-point units would also help.

Floating-Point Unit Organization



Alternate, Faster Floating-Point Unit Organization



The current fp organization doesn't make good use of pipelining; one operation must complete before the next one can start. A faster fp organization allowing the use of pipelining is shown to the right. Other organizations with good performance are possible. The current unit focuses on small size.

Note that the ISA isn't closely related to the implementation of the floating-point unit. There is no reason why faster floating point wouldn't be possible from an ISA perspective.

Flow Control Unit / Branch Unit

A single flow control (or branch) unit takes care of all the flow control instructions the core supports. The author prefers to call the unit a flow control unit rather than a branch unit, because the unit takes care of additional instructions besides simple branches. A branch implies multiple paths of execution. For some instructions for instance a jump or a call there is only one path of execution, calling them a branch is a bit of a misnomer. In an initial version of the core flow control operations were performed by the alu's. This led to problems of prioritization of flow control operations when two flow control operations were taking place at the same time. Especially given that the operations may have been assigned out of order to the alu's. Moving from the alu pair to a single flow control unit resolved those problems.

The flow control unit is responsible for determining whether a branch should be taken. However, by the time the branch reaches the flow control unit, it has already taken a predicted path of execution. So, the flow control unit's real job is to verify that the correct path was taken.

The flow control unit has a small alu to calculate register increment or decrement, stack pointer adjustment, and return addresses. The fcu's alu is called NVIO_fcu_calc.v to avoid confusion with the master alu.

Memory Unit

The memory unit can handle up to two requests at a time. If the data cache is enabled loads check for data in the data cache, which is loaded with data if the data is not present in the cache. Up to two load operations may be taking place at the same time, each one making use of a different read port of the data cache.

Un-cached loads and stores access external memory and hence are serialized. Since there is only a single port to external memory access takes place one request at a time.

Stores

The write buffer is dual ported and will accept two store operations at the same time. They will be placed in the write queue in program order.

Large immediates

Occasionally there is a need for large constants. These can be built up using shifting instructions in the instruction set. Or, a large constant may be supplied directly in the instruction stream via the use of the large constant templates. An 80-bit constant occupies slots 1 and 2 for these templates. It does mean that it is necessary to align the instruction appropriately which may require adding NOP instructions into the instruction stream. However, it may require up to five equivalent instructions to build up the immediate constant using other instructions. It also uses at least one additional register.

The core will treat instruction bundles that contain large immediate (templates 7D, 7E) as a single instruction. There needs to be only a single queue slot available for the core to queue one of these instructions.

Operating Levels

The core has four operating levels. The highest operating level is operating level zero which is called the machine operating level. Operating level zero has complete access to the machine. Other operating levels may have more restricted access. When an interrupt occurs, the operating level is set to the machine level. The core vectors to an address depending on the current operating level. When operating at level zero addresses are not subjected to translation and the virtual address and physical address are the same.

Operating Level	Privilege Level	Moniker
3	7 to 255	user
2	2 to 6	supervisor
1	1	hypervisor
0	0	machine

Switching Operating Levels

The operating level is automatically switched to the machine level when an interrupt occurs. The BRK instruction may be used to switch operating levels. The REX instruction may also be used by an interrupt handler to switch the operating level to a lower level. The RTI instruction will switch the operating level back to what it was prior to the interrupt.

Privilege Levels

The core supports a 256-level privilege level system. Privilege level zero is assigned to operating level zero. Privilege level one is assigned to operating level one. Privilege levels 2 to 6 are assigned to operating level two. The remaining privilege levels are assigned to operating level three.

Debug Mode

In debug mode the core single steps one instruction at a time.

TLB – The Translation Lookaside Buffer

Overview

The TLB (translation look-aside buffer) offers a means of address virtualization and memory protection. A TLB works by caching address mappings between a real physical address and a virtual address used by software. The TLB is managed by software triggered when a TLB miss occurs. The TLB deals with memory organized as pages. Typically, software manages a paging table whose entries are loaded into the TLB as translations are required.

The TLB keeps a reference count for each map entry stored in the TLB. The upper 24-bits of the reference count, which is a 32-bit saturating counter, are automatically incremented with each memory access to the page. Reference counts are subject to aging under control of the AFC register. The reference counts may be read or written with the TBLRDAGE or TLBWRAGE commands.

The TLB is manipulated with the [TLB](#) instruction.

Size / Organization


The core uses a 256 entry TLB (translation look-aside buffer) to support virtual memory. The TLB supports variable page sizes from 8kB to 2MB. The TLB is organized as a sixteen-way sixteen-set cache. The TLB processes all addresses leaving the core including both code and data addresses.

Updating the TLB

The TLB is updated by first placing values into the TLB holding registers using the TLB instruction, then issuing a TLB write command using the TLB command instruction.

Address translations will not take place until the TLB is enabled. An enable TLB command must be issued using the TLB command instruction.

TLB Entries



G = Global

The global bit marks the TLB entry as a global address translation where the ASID field is not used to match addresses.

ASID = address space identifier

The ASID field in the TLB entry must match the processor's current ASID value for the translation to be considered valid, unless the G bit is set. If the G bit is set in the TLB

entry, then the ASID field is ignored during the address comparison. The processor's current ASID is located in the machine status register.

C = cache-ability bits

If the cache-ability bits are set to 001_b then the page is un-cached, otherwise the page is cached.

D = dirty bit

The dirty bit is set by hardware when a write occurs to the virtual memory page identified by the TLB entry.

A = accessed bit

This bit is set when the page is accessed.

U = undefined usage

This bit is available for OS usage

S = address shortcut

R = read bit

This bit indicates that the page is readable.

W = write bit

This bit indicates that the page is writeable

X = execute bit

This bit indicates that the page contains executable code

R,W,X = valid bit

One of these bits must be set for the address translation to be considered valid. The entire TLB may be invalidated using the invalidate all command.

Page Table Entry

The following layout shows the page table entry structure as stored in memory. Although the page table is managed by software, this layout should be followed.

31			16	15		8	7	6	5	4	3	2	1	0
Share Count ₁₆				PrivLevel		D	U	S	A	C	R	W	X	
ASID ₈		~ ₄		Protection Key ₂₀										
Reference Counter ₃₂														
Page Number ₃₂														

Bit			
0	X	1 = executable	Together these three fields combined indicate if the page is present. It must be at least one of readable, writeable, or executable.
1	W	1 = page writeable	
2	R	1 = readable	
3	C	1 = cachable	Ignored for executable pages which are always cached
4	A	1 = accessed	
5	S	1 = shortcut translation	Translation shortcut bit eg (8MiB pages)
6	U	undefined usage	available to be used by OS
7	D	1=dirty	set if the page is written to
8 to 15	PL	Privilege level	
16 to 31	SC	Shared memory reference count	Number of times the memory page is shared.
32 to 51	KY	Protection key	20-bit key used to protect memory page
52 to 55	Pad1	reserved area	
56 to 63	ASID	Address Space IDentifier	
64 to 95	RC	Reference Count	32-bit page reference count
96 to 127	PN	Memory page number	virtual page number associated with physical page

Page Directory Entry

Page directory entries have the same format as page table entries.

TLB Registers

TLBWired (#0h)

This register limits random updates to the TLB to a subset of the available number of ways. TLB ways below the value specified in the Wired register will not be updated randomly. Setting this register provides a means to create fixed translation settings. For instance, if the wired register is set to two, the thirty-two fixed entries will be available.

TLBIndex (#1h)

This register contains the entry number of the TLB entry to be read from or written to.

TLBRandom (#2h)

This register contains a random four-bit value used to update a random TLB entry during a TLB write operation.

TLBPageSize (#3h)

The TLBPageSize register controls which address bits are significant during a TLB lookup.

N	Page Size	
0	8KiB	
1	32kiB	
2	128kiB	
3	512kiB	
4	2MiB	
5	8MiB	

TLBPhysPage (#5h)

The TLBPhysPage register is a holding register that contains the page number for an associated virtual address. This register is transferred to or from the TLB by TLB instructions.

79	0
Physical Page Number	

TLBVirtPage (#4h)

The TLBVirtPage register is a holding register that contains the page number for an associated physical address. This register is transferred to or from the TLB by TLB instructions.

79	0
Virtual Page Number	

TLBASID (#7h)

The TLBASID register is a holding register that contains the address space identifier (ASID) , valid, dirty, global, and cache-ability bits associated with a TLB entry. This register is transferred to or from the TLB by TLB instructions.

79	32	31	24	23	16	13	11	10	9	8	7	6	5	3	2	1	0
-----	PL			ASID		PgSz		G	D	U	S	A	C	R	W	X	

TLBAFC (#12) – Aging Frequency Control

This 24-bit register controls the frequency of aging applied to page reference counters. The aging counter is decremented at the core's clock frequency. When the counter underflows it triggers an aging cycle and is reloaded with the value in the AFC register. The AFC register is defaulted to 20000. This gives an aging frequency of 1000Hz with a 20MHz core clock.

Exceptions

External Interrupts

There is very little difference between an externally generated exception and an internally generated one. An externally caused exception will force a BRK instruction into the instruction stream. The BRK instruction contains a cause code identifying the external interrupt source.

Polling for Interrupts

To support managed code an interrupt polling instruction (PFI) is provided in the instruction set. In some managed code environments, it is not enough to disable and enable interrupts around critical code. The code must be effectively run with interrupt disabled all the time. This makes it necessary to poll for interrupts in software. For instance, stack prologue code may cause false pointer matches for the garbage collector because stack space is allocated before the contents are defined. If the GC scan occurs on this allocated but undefined area of memory, there could be false matches.

Effect on Machine Status

The operating mode is always switched to the machine mode on exception. It's up to the machine mode code to redirect the exception to a lower operating mode when desired. Further exceptions at the same or lower interrupt level are disabled automatically. Machine mode code must enable interrupts at some point. This can be done automatically when the exception is redirected to a lower level by the REX instruction. The RTI instruction will also automatically enable further machine level exceptions.

Exception Stack

The instruction pointer and status bits are pushed onto an internal stack when an exception occurs. This stack is only eight entries deep as that is the maximum amount of nesting that can occur. Further nesting of exceptions can be achieved by saving the state contained in the exception registers.

Exception Vectoring

Exceptions are handled through a vector table. The vector table has four entries, one for each operating level the core may be running at. The location of the vector table is determined by TVEC[0]. If the core is operating at level three for instance and an interrupt occurs vector table address number three is used for the interrupt handler. Note that the interrupt automatically switches the core to operating level zero, privilege level zero. An exception handler at the machine level may redirect exceptions to a lower level handler identified in one of the vector registers. More specific exception information is supplied in the cause register.

Operating Level	Address (If TVEC[0] contains \$FFFC0000)	
0	\$FFFC0000	Handler for operating level zero
1	\$FFFC0020	
2	\$FFFC0040	
3	\$FFFC0060	

Reset

The core begins executing instructions at address \$FFFFFFFFFFFFFFFC0100. All registers are in an undefined state. Register set #0 is selected.

Precision

Exceptions in NVIO are precise. They are processed according to program order of the instructions. If an exception occurs during the execution of an instruction, then an exception field is set in the reorder buffer. The exception is processed when the instruction commits which happens in program order. If the instruction was executed in a speculative fashion, then no exception processing will be invoked unless the instruction makes it to the commit stage.

Exception Cause Codes

The following table outlines the cause code for a given purpose. These codes are specific to NVIO. Under the HW column an 'x' indicates that the exception is internally generated by the processor; the cause code is hard-wired to that use. An 'e' indicates an externally generated interrupt, the usage may vary depending on the system.

Cause Code		HW	Description	
1	IBE	x	instruction bus error	
2	EXF	x	Executable fault	
4	TLB	x	tlb miss	
			FMTK Scheduler	
128		e		
129	KRST	e	Keyboard reset interrupt	
130	MSI	e	Millisecond Interrupt	
131	TICK	e		
156	KBD	e	Keyboard interrupt	
157	GCS	e	Garbage collect stop	
158	GC	e	Garbage collect	
159	TSI	e	FMTK Time Slice Interrupt	
3			Control-C pressed	
20			Control-T pressed	
26			Control-Z pressed	
32	SSM	x	single step	
33	DBG	x	debug exception	
34	TGT	x	call target exception	
35	MEM	x	memory fault	
36	IADR	x	bad instruction address	
37	UNIMP	x	unimplemented instruction	
38	FLT	x	floating point exception	
39	CHK	x	bounds check exception	
40	DBZ	x	divide by zero	
41	OFL	x	overflow	
47				
48	ALN	x	data alignment	
49				
50	DWF	x	Data write fault	
51	DRF	x	data read fault	
52	SGB	x	segment bounds violation	
53	PRIV	x	privilege level violation	
54	CMT	x	commit timeout	
55	BD	x	branch displacement	
56	STK	x	stack fault	
57	CPF	x	code page fault	
58	DPF	x	data page fault	

60	DBE	x	data bus error	
61				
62	NMI	x	Non-maskable interrupt	
225	FPX_IOP	x	Floating point invalid operation	
226	FPX_DBZ	x	Floating point divide by zero	
227	FPX_OVER	x	floating point overflow	
228	FPX_UNDER	x	floating point underflow	
229	FPX_INEXACT	x	floating point inexact	
231	FPX_SWT	x	floating point software triggered	
240	SYS		Call operating system (FMTK)	
241			FMTK Schedule interrupt	
255	PFI		reserved for poll-for-interrupt instruction	

DBG

A debug exception occurs if there is a match between a data or instruction address and an address in one of the debug address registers.

IADR

This exception is currently not implemented but reserved for the purpose of identifying bad instruction addresses. If the two least significant bits of the instruction address are non-zero then this exception will occur.

UNIMP

This exception occurs if an instruction is encountered that is not supported by the processor. It is not currently implemented.

OFL

If an arithmetic operation overflows (multiply, add, or shift) and the overflow exception is enabled in the arithmetic exception enable register then an OFL exception will be triggered.

FLT

A floating-point exception is triggered if an exceptional condition occurs in the floating-point unit and the exception is enabled. Please see the section on floating-point for more details.

DRF, DWF, EXF

Data read fault, data write fault, and execute fault are exceptions that are returned by the memory management unit when an attempt is made to access memory for which the corresponding access type is not allowed. For instance, if the memory page is marked as non-executable an attempt is made to load the instruction cache from the page then an execute fault EXF exception will occur.

CPF, DPF

The code page fault and data page fault exceptions are activated by the mmu if the page is not present in memory. Access may be allowed but simply unavailable.

PRIV

Some instructions and CSR registers are legal to use only at a higher operating level. If an attempt is made to use the privileged instruction by a lower operating level, then a privilege violation exception may occur. For instance, attempting to use RTI instruction from user operating level.

STK

If the value loaded into one of the stack pointer registers (the stack point sp or frame pointer fp) is outside of the bounds defined by the stack bounds registers, then a stack fault exception will be triggered.

DBE

A timeout signal is typically wired to the err_i input of the core and if the data memory does not respond with an ack_i signal fast enough an error will be triggered. This will happen most often when the core is attempting to access an unimplemented memory area for which no ack signal is generated. When the err_i input is activated during a data fetch, an exception is flagged in a result register for the instruction. The core will process the exception when the instruction commits. If the instruction does not commit (it could be a speculated load instruction) then the exception will not be processed.

IBE

A timeout signal is typically wired to the err_i input of the core and if the instruction memory does not respond with an ack_i signal fast enough an error will be triggered. This will happen most often when the core is attempting to access an unimplemented memory area for which no ack signal is generated. When the err_i input is activated during an instruction fetch, a breakpoint instruction is loaded into the cache at the address of the error.

NMI

The core does not currently support non-maskable interrupts. However, this cause value is reserved for that purpose.

BD

The core will generate the BD (branch displacement) exception if a branch instruction points back to itself. Branch instructions in this sense include jump (JMP) and call (CALL) instructions.

Bundle Format:

127	120	119	80	79	40	39	0
Template ₈		Slot2			Slot1		Slot0

Instruction Formats:

Operation Sizes

Many instructions have an option to process data in sub-word data sizes including bytes, wydes, and penta-bytes. Typically, sized operations are supported only with register-register instructions. Instructions using immediate values always operate on whole words.

SIMD

Single instruction multiple data operations treat the 80-bit operands as multiple independent lanes of data depending on the size selected. For a half-word size, the operands are treated as two independent 40-bit operands. For a wyde size the operands are treated as five independent 16-bit operands. SIMD operations are selected by setting the parallel operation bit in the instruction (the most significant bit of the size field).

Arithmetic Operations

Arithmetic operations include addition, subtraction, comparison, multiplication and division.

Relational Operations

There are handful of instructions for determining the relationship between values. These include an assortment of set instructions.

Logical Operations

Logical operations include bitwise and, or, and exclusive or. Inverted logical ops are also available for register instruction forms (nand, nor, and exnor).

Shift Operations

There is a full complement of shift operations including left and right unsigned and signed shifts and left and right rotate instructions.

Bitfield Operations

The CC80 compiler has direct support for bitfields and bitfield instructions support these operations. Bitfield operations include insert, extract, set, change and clear operations.

Memory Operations

Memory operations include loads and stores of bytes, wydes or half-words. The core can perform loads and stores using indexed addressing.

Instruction Formats:

Imm ₅	Opc ₂	Imm ₁₇					Rs1 ₆	Opcode ₄	Rd ₆	ML					
Imm ₅	Opc ₂	Imm ₁₁				Rs2 ₆		Rs1 ₆	Opcode ₄	Imm ₆	MS				
Imm ₅	2 ₂	Imm ₁₁				DC ₃	IC ₃	Rs1 ₆	Eh ₄	Imm ₆	CACHE				
0Ah ₅	2 ₂	~	~	~ ₃	~ ₆	Rs2 ₆		63 ₆	Dh ₄	63 ₆	PUSH				
Imm ₂₄								63 ₆	Bh ₄	63 ₆	PUSHC				
0Ah ₅	1 ₂	~	~	~ ₃	~ ₆	~ ₆		63 ₆	Dh ₄	Rd ₆	POP				
Funct ₅	0 ₂	~	~	Sc ₃	Rs3 ₆	Imm ₄	Md ₂	Rs1 ₆	Fh ₄	Rd ₆	MLX				
Funct ₅	2 ₂	~	~	Sc ₃	Rs3 ₆	Rs2 ₆		Rs1 ₆	Fh ₄	Imm ₄	Md ₂	MSX			
Funct ₅	0 ₂	A	R	Sz ₃	~ ₆	Rs2 ₆		Rs1 ₆	Ah ₄	Rd ₆	AMO				
Funct ₅	0 ₂	A	R	Sz ₃	Imm ₁₂			Rs1 ₆	Ah ₄	Rd ₆	AMOI				
0Eh ₅	2 ₂	~	~	Sc ₃	Rs3 ₆	DC ₃	IC ₃	Rs1 ₆	Fh ₄	Imm ₆	CACHEX				
19h ₅	2 ₂	~	~	~ ₃	~ ₆	~ ₆		~ ₆	Fh ₄	~ ₆	MEMDB				
18h ₅	2 ₂	~	~	~ ₃	~ ₆	~ ₆		~ ₆	Fh ₄	~ ₆	MEMSB				
Cmd ₅	2 ₂	~	~	~ ₃	~ ₆	~ ₂	Tn ₄	Rs1 ₆	Ch ₄	Rd ₆	TLB				
Imm ₅	3 ₂	Imm ₁₁				62 ₆		63 ₆	Fh ₄	62 ₆	LINK				
0Ah ₅	1 ₂	~	~	~ ₉		62 ₆		63 ₆	Fh ₄	62 ₆	UNLK				
Imm ₅	Op ₂	Imm ₁₇						Rs1 ₆	Opcode ₄	Rd ₆	RI				
Funct ₅	0 ₂	0 ₂		~ ₃	~ ₆	~ ₆		Rs1 ₆	1 ₄	Rd ₆	R1				
Funct ₅	Op ₂	Fn ₂		Sz ₃	Rs3 ₆	Rs2 ₆		Rs1 ₆	Opcode ₄	Rd ₆	R3				
Funct ₅	Op ₂	Fn ₂		Sz ₃	Rs3 ₆	Rs2 ₆		Rs1 ₆	Opcode ₄	Rd ₆	SHIFT				
Funct ₅	Op ₂	Fn ₂		Sz ₃	Rs3 ₆	Imm ₆		Rs1 ₆	Opcode ₄	Rd ₆	SHIFTI				
Funct ₄	Rg ₃	o	w	~ ₃	Rs3 ₆	Bw ₆		Bo ₆	Bh ₄	Rd ₆	BF				
Funct ₄	Rg ₃	o	Bw ₄		Rs3 ₆	Rs2 ₆		Bo ₆	Bh ₄	Rd ₆	BFINS				
Op ₂	OL ₂	~ ₃		0 ₂	~ ₃	Regno ₁₂		Rs1 ₆	5 ₄	Rd ₆	CSR				
Displacement _{19..12}		Offset _{11..4}			S ₂	Rs2 ₆		Rs1 ₆	Opcode ₄	P ₂	~	Cond ₃	Bcc		
Displacement _{19..12}		Offset _{11..4}			S ₂	Bitno _{6..1}		Rs1 ₆	5h ₄	P ₂	~	B ₀	Cn ₂	BBS	
Displacement _{19..12}		Offset _{11..4}			S ₂	Imm ₆		Rs1 ₆	Opcode ₄	P ₂	~	Imm ₃		BEQI	
~ ₉				~ ₃	Rs3 ₆		Rs2 ₆		Rs1 ₆	2h ₄		~ ₂	Cond ₄		BRG
~ ₉				~ ₃	~ ₆		~ ₆		~ ₆	3h ₄		~ ₆		NOP	
Imm ₂₄								Rs1 ₆	8h ₄	61 ₆		JAL			
Address _{37..8}									Opcode ₄	Address _{7..2}			CALL		
Imm _{19..1}						61 ₆		63 ₆	Bh ₄	63 ₆		RET			
H	0 ₄	~ ₅			Cause ₈		~ ₂	Imask ₄	Rs1 ₆	Fh ₄	~ ₆		BRK		
0	1 ₄	~ ₅			255 ₈		~ ₂	0 ₄	0 ₆	Fh ₄	~ ₆		PFI		

0 ₅	~7		~6		~6		Rs1 ₆	Eh ₄	Sema ₆	RTI	
1 ₅	~5		PrivLvl ₈		T ₂	Imask ₄	Rs1 ₆	Eh ₄	~6	REX	
2 ₅	~7			~6		~6		~6	Eh ₄	~6	SYNC
3 ₅	~7			~6		~2	Imask ₄	Rs1 ₆	Eh ₄	Rd ₆	SEI
4 ₅	~7			Imm ₆		Rs2 ₆		Rs1 ₆	Eh ₄	~6	WAIT
5 ₅	~7			~6		~6		Codebuf ₆	Eh ₄	~6	EXEC
Imm ₁₈						Rs2 ₆		Rs1 ₆	Ch ₄	Imm ₆	CHKI
~12				Rs3 ₆		Rs2 ₆		Rs1 ₆	Dh ₄	~6	CHK
Funct ₅	Prec ₄		Rm ₃	Rs3 ₆		Rs2 ₆		Rs1 ₆	3 ₄	Rd ₆	FLT3
Funct ₅	Prec ₄		Rm ₃	~6		Rs2 ₆		Rs1 ₆	2 ₄	Rd ₆	FLT2
Funct ₅	Prec ₄		Rm ₃	~2	Op ₃	Rs2 ₆		Rs1 ₆	2 ₄	Rd ₆	FSGNOP
Funct ₅	Prec ₄		Rm ₃	~6		Cndx ₆		Rs1 ₆	6 ₄	Rd ₆	FLT2I
Funct ₅	Prec ₄		Rm ₃	~6		~6		Rs1 ₆	1 ₄	Rd ₆	FLT1
Funct ₅	Prec ₄		Rm ₃	~6		~6		Rs1 ₆	A h ₄	Rd ₆	FLTLI
Funct ₅	Prec ₄		Rm ₃	~6		Imm ₆		Rs1 ₆	1 ₄	Rd ₆	FXX
Imm ₇		Op ₂	Imm ₁₅				Rs1 ₆		Opcode ₄	Rd ₆	FLTLDI

Template bits 0 to 6 combined with the instruction slot number determine which functional unit the instructions is for. Template bit 7 is reserved and should be set to zero.

Instruction Encodings

Memory Operations {Opc2, Opcode4}

	x0	x1	x2	x3	x4	x5	x6	x7	x8	x9	xA	xB	xC	xD	xE	xF
0x	LDB	LDW	LDP	LDD	LDBU	LDWU	LDPU	LDDR	LDT	LDO	{AMO}	PUSHC	LDTU	LDOU	LEA	{MLX}
1x	LDFS	LDfD	LDfDP	LDDP								PUSHC		POP	NOP	UNLK
2x	STB	STW	STP	STD				STDC	STT	STO	CAS	PUSHC	TLB	PUSH	CACHE	{MSX}
3x	STFS	STFD	STfDP	STDP								PUSHC			NOP	LINK

Indexed Memory Loads {MLX} {Funct5}

Funct ₅	xx000	xx001	xx010	xx011	xx100	xx101	xx110	xx111
00xxx	LDBX	LDWX	LDPX	LDDX	LDBUX	LDWUX	LDPUX	LDDRX
01xxx	LDTX	LDOX			LDTUX	LDOUX	LEAX	
10xxx								
11xxx			LMX					

AMO {Funct5}

Funct ₅	xx000	xx001	xx010	xx011	xx100	xx101	xx110	xx111
00xxx	SWAP	SWAPI	ADD	ADDI	AND	ANDI	OR	ORI
01xxx	XOR	XORI	SHL	SHLI	SHR	SHRI	MIN	MINI
10xxx	MAX	MAXI	MINU	MINUI	MAXU	MAXUI		
11xxx								

Indexed Memory Stores / Miscellaneous {Funct5}

Funct ₅	xx000	xx001	xx010	xx011	xx100	xx101	xx110	xx111
00xxx	STBX	STWX	STPX	STDx				STDCX
01xxx	STTX	STOX	CASX				CACHE	
10xxx								
11xxx	MEMSB	MEMDB	SMX					

Flow Control

Opcode ₄	xx00	xx01	xx10	xx11
00xx	Bcc	BLcc	BRcc	NOP
01xx	FBcc	BBc	BEQ #	BNE #
10xx	JAL	JMP	CALL	RET
11xx	CHK #	CHK	RTI / REX / Misc	BRK

Branch Predictor Bits

P ₂	Prediction
0	No prediction – use branch predictor
1	reserved
2	predict as not taken
3	predict as taken

Unit		
0		
1	Branch	
2	Integer	
3	Floating Point	
4	Memory	
5		
6		
7		

Integer ALU {bits 32, 31, Opcode₄}

	x0	x1	x2	x3	x4	x5	x6	x7	x8	x9	xA	xB	xC	xD	xE	xF
0x		{R1}	{R3}	{R3}	ADD	CSR	CMP	CMPU	AND	OR	XOR	{BF}	BLEND			MULF
1x	SLT	SGE	SLE	SGT	SLTU	SGEU	SLEU	SGTU	SEQ	SNE		{BF}	BYTNDX	WYDNDX		DIF
2x	MUL	MULU	DIV	DIVU	MOD	MODU			MADF			{BF}				
3x	FXMUL			ADDS0	ADDS1	ADDS2	ADDS3	ANDS0	ANDS1	ANDS2	ANDS3	{BF}	ORS1	ORS2	ORS3	ORS0

ALU – R3 Format {Funct₅, bit 6}

	x0	x1	x2	x3	x4	x5	x6	x7	x8	x9	xA	xB	xC	xD	xE	xF
0x	MULH	MULUH	ADDV	SUBV	ADD	SUB	CMP	CMPU	AND	OR	XOR		NAND	NOR	XNOR	MULF
1x	SLT	SGE	SLE	SGT	SLTU	SGEU	SLEU	SGTU	SEQ	SNE		CMOVNZ	MIN	MAX	PTRDIF	DIF
2x	MUL	MULU	DIV	DIVU	MOD	MODU			MADF	MUX	BYTNDX	WYDNDX	MAJ	AVG		
3x	FXMUL	FXDIV	SHL	ASL	SHR	ASR	ROL	ROR	SHL #	ASL #	SHR #	ASR #	ROL #	ROR #	BMM	

ALU – R1 Format Funct₅

	x0	x1	x2	x3	x4	x5	x6	x7	x8	x9	xA	xB	xC	xD	xE	xF
0x	CNTLZ	CNTLO	CNTPOP	COM	ABS	NOT	ISPTR	NEG	ZXT	ZXC	ZXB	ZXP	ZXO			
1x	MOV								SXT	SXC	SXB	SXP	SXO			

Instruction Descriptions

Execution Units: this is the functional unit that the instruction is executed by.

Clock cycles:

This is a highly optimistic estimate of the number of clock cycles required to complete the operation. This estimate assumes all values required for the operation are available without delay. Under ideal circumstances many operations can complete in 1/3 of a clock cycle because there are enough functional units to service that many instructions at once.

Branch Unit

The branch unit executes instructions that control the flow of program execution. These instructions include things like SEI which enables interrupts in addition to regular branch, jump and call instructions.

Branch Targets

The core uses bundle relative addressing. The target slot within a bundle is encoded separately (The S_2 field of the instruction.)

Branch Predictions

Normally the branch predictor determines the branch prediction, however, a branch instruction may be statically predicted to be taken or not taken. This is controlled by the P_2 field of the branch instruction.

BAND –Branch on Logical And

Description:

If the logical and of two registers is true, a target address is loaded into the instruction pointer. If the branch branches back to itself a branch exception will be generated. The branch range is approximately +/-2MB. Non-zero values are considered to be true and zero is considered to be false.

Instruction Format: Bcc, BRcc

Operation:

```
if (Rs1 && Rs2)
    ip = target
```

Clock Cycles: 0.33 (Ideal) 1 (average)

Execution Units: Branch

Exceptions: branch target

BBC –Branch if Bit Clear

Description:

If the specified bit in a register is clear, the target address is loaded into the instruction pointer. If the branch branches back to itself a branch exception will be generated. The range of the branch is approximately 2MB.

Instruction Format: BBc

Operation:

if (Ra[bitno]=0)
ip = target

Clock Cycles: 0.33 (Ideal) 1 (average)

Execution Units: branch

Exceptions: branch target address

BBS –Branch if Bit Set

Description:

If the specified bit in a register is set, the target address is loaded into the instruction pointer. If the branch branches back to itself a branch exception will be generated. The range of the branch is approximately 2MB.

Instruction Format: BBc

Operation:

if (Ra[bitno]=1)
ip = target

Clock Cycles: 0.33 (Ideal) 1 (average)

Execution Units: branch

Exceptions: branch target address

Bcc – Conditional Branch

Description:

If the branch condition is true, the target address is loaded into the instruction pointer. If the instruction branches back to itself, a branch target exception will occur. The range of the branch is approximately 2MB.

Instruction Format: Bcc

Cond ₃	Mne.	
0	BEQ	Rs1 = Rs2 signed
1	BNE	Rs1 \neq Rs2
2	BLT / BGT	Rs1 < Rs2 or Rs2 > Rs1
3	BGE / BLE	Rs1 \geq Rs2 or Rs2 \leq Rs1
4	reserved	
5	reserved	
6	BLTU	Rs1 < Rs2 (unsigned)
7	BGEU	Rs1 \geq Rs2 (unsigned)

Clock Cycles: 0.33 (Ideal) 1 (average)

Execution Units: Branch

Exceptions: branch target

BEQ –Branch if Equal

Description:

If the two registers are equal, a target address is loaded into the instruction pointer. If the branch branches back to itself a branch exception will be generated. The branch range is approximately 2MB.

Instruction Format: Bcc, BRcc

Operation:

if ($Rs1 = Rs2$)
 $ip[22:0] = target$

Clock Cycles: 0.33 (Ideal) 1 (average)

Execution Units: Branch

Exceptions: branch target

BEQI –Branch if Equal Immediate

Description:

If a register is equal to nine-bit sign extended value, a target address is loaded into the instruction pointer. If the branch branches back to itself a branch exception will be generated. The branch range is approximately 2MB. This instruction is particularly useful for implementing case statements based on small values.

Instruction Format: Bcc

There is no register target form for this instruction.

Operation:

if ($Ra = Imm_9$)
ip = target

Clock Cycles: 0.33 (Ideal) 1 (average)

Execution Units: Branch

Exceptions: branch target

BGE –Branch if Greater or Equal

Description:

If the value in Rs1 is greater than or equal to the value in Rs2, a target address is loaded into the instruction pointer. The values are treated as signed numbers. If the branch branches back to itself a branch exception will be generated. The branch range is approximately 2MB.

Instruction Format: Bcc, BRcc

Operation:

if (Rs1 \geq Rs2)
ip = target

Clock Cycles: 0.33 (Ideal) 1 (average)

Execution Units: Branch

Exceptions: branch target

BGEU –Branch if Greater or Equal

Description:

If the value in Rs1 is greater than or equal to the value in Rs2, a target address is loaded into the instruction pointer. The values are treated as unsigned numbers. If the branch branches back to itself a branch exception will be generated. The branch range is approximately 2MB.

Instruction Format: Bcc, BRcc

Operation:

if (Rs1 \geq Rs2)
ip = target

Clock Cycles: 0.33 (Ideal) 1 (average)

Execution Units: Branch

Exceptions: branch target

BGT –Branch if Greater Than

Description:

If the value in Rs2 is greater than the value in Rs1, a target address is loaded into the instruction pointer. The values are treated as signed numbers. If the branch branches back to itself a branch exception will be generated. The branch range is approximately 2MB. This is an alternate mnemonic for the BLT instruction. The register fields are swapped.

Instruction Format: Bcc, BRcc

Operation:

if (Rs2 > Rs1)
ip = target

Clock Cycles: 0.33 (Ideal) 1 (average)

Execution Units: Branch

Exceptions: branch target

BGTU –Branch if Greater Than

Description:

If the value in Rs2 is greater than the value in Rs1, a target address is loaded into the instruction pointer. The values are treated as unsigned numbers. If the branch branches back to itself a branch exception will be generated. The branch range is approximately 2MB. This is an alternate mnemonic for the BLTU instruction.

Instruction Format: Bcc, BRcc

Operation:

if (Rs2 > Rs1)
ip = target

Clock Cycles: 0.33 (Ideal) 1 (average)

Execution Units: Branch

Exceptions: branch target

BLE –Branch if Less Than or Equal

Description:

If the value in Rs2 is less than or equal to the value in Rs1, a target address is loaded into the instruction pointer. The values are treated as signed numbers. If the branch branches back to itself a branch exception will be generated. The branch range is approximately 2MB. This is an alternate mnemonic for the BGE instruction.

Instruction Format: Bcc, BRcc

Operation:

```
if (Rs2 <= Rs1)
    ip = target
```

Clock Cycles: 0.33 (Ideal) 1 (average)

Execution Units: Branch

Exceptions: branch target

BLEU –Branch if Less Than or Equal

Description:

If the value in Rs2 is less than or equal to the value in Rs1, a target address is loaded into the instruction pointer. The values are treated as unsigned numbers. If the branch branches back to itself a branch exception will be generated. The branch range is approximately 2MB. This is an alternate mnemonic for the BGEU instruction.

Instruction Format: Bcc, BRcc

Operation:

if (Rs2 <= Rs1)
ip = target

Clock Cycles: 0.33 (Ideal) 1 (average)

Execution Units: Branch

Exceptions: branch target

BLT –Branch if Less Than

Description:

If the value in Rs1 is less than the value in Rs2, a target address is loaded into the instruction pointer. The values are treated as signed numbers. If the branch branches back to itself a branch exception will be generated. The branch range is approximately 2MB.

Instruction Format: Bcc, BRcc

Operation:

if (Rs1 < Rs2)
ip = target

Clock Cycles: 0.33 (Ideal) 1 (average)

Execution Units: Branch

Exceptions: branch target

BLTU –Branch if Less Than

Description:

If the value in Rs1 is less than the value in Rs2, a target address is loaded into the instruction pointer. The values are treated as unsigned numbers. If the branch branches back to itself a branch exception will be generated. The branch range is approximately 2MB.

Instruction Format: Bcc, BRcc

Operation:

if (Rs1 < Rs2)
ip = target

Clock Cycles: 0.33 (Ideal) 1 (average)

Execution Units: Branch

Exceptions: branch target

BNAND –Branch on Logical Nand

Description:

If the logical nand of two registers is true, a target address is loaded into the instruction pointer. If the branch branches back to itself a branch exception will be generated. The branch range is approximately 2MB. Non-zero values are considered to be true and zero is considered to be false.

Instruction Format: Bcc, BRcc

Operation:

```
if (!(Ra && Rb))  
    ip = target
```

Clock Cycles: 0.33 (Ideal) 1 (average)

Execution Units: Branch

Exceptions: branch target

BNE –Branch if Not Equal

Description:

If the two registers are not equal, a target address is loaded into the instruction pointer. If the branch branches back to itself a branch exception will be generated. The branch range is approximately 2MB.

Instruction Format: Bcc, BRcc

Operation:

if (Rs1 \neq Rs2)
ip = target

Clock Cycles: 0.33 (Ideal) 1 (average)

Execution Units: Branch

Exceptions: branch target

BNOR –Branch on Logical Nor

Description:

If the logical or of two registers is false, a target address is loaded into the instruction pointer. If the branch branches back to itself a branch exception will be generated. The branch range is approximately 2MB. Non-zero values are considered to be true and zero is considered to be false.

Instruction Format: Bcc, BRcc

Operation:

```
if (!(Rs1 || Rs2))  
    ip = target
```

Clock Cycles: 0.33 (Ideal) 1 (average)

Execution Units: Branch

Exceptions: branch target

BOR –Branch on Logical Or

Description:

If the logical or of two registers is true, a target address is loaded into the instruction pointer. If the branch branches back to itself a branch exception will be generated. The branch range is approximately 2MB. Non-zero values are considered to be true and zero is considered to be false.

Instruction Format: Bcc, BRcc

Operation:

if (Rs1 || Rs2)
ip = target

Clock Cycles: 0.33 (Ideal) 1 (average)

Execution Units: Branch

Exceptions: branch target

BRcc – Conditional Branch to Register

Description:

This instruction allows the target address to be supplied by a register, accommodating computed target addresses. If the branch condition is true, the target address is loaded into the instruction pointer. All bits of the instruction pointer are affected. If the instruction branches back to itself, a branch target exception will occur.

Instruction Format: BRcc

Cond ₄	Mne.	
0	BEQ	Rs1 = Rs2 signed
1	BNE	Rs1 \neq Rs2
2	BLT / BGT	Rs1 < Rs2 or Rs2 > Rs1
3	BGE / BLE	Rs1 >= Rs2 or Rs2 <= Rs1
4	BNAND	!(Rs1 && Rs2)
5	BNOR	!(Rs1 Rs2)
6	BLTU	Rs1 < Rs2 (unsigned)
7	BGEU	Rs1 >= Rs2 (unsigned)
8	FBEQ	Rs1 = Rs2 floating-point compare
9	FBNE	Rs1 \neq Rs2
A	FBLT	Rs1 < Rs2
B	FBLE	Rs1 <= Rs2
C	BAND	Rs1 && Rs2
D	BOR	Rs1 Rs2
E	reserved	
F	FBUN	unordered comparison

Clock Cycles:

Execution Units: Branch

Exceptions: branch target

BRK – Hardware / Software Breakpoint

Description:

Invoke the break handler routine. The break handler routine handles all the hardware and software exceptions in the core. A cause code is loaded into the CAUSE CSR register. The break handler should read the CAUSE code to determine what to do. The break handler is located by TVEC[0]. This address should contain a jump to the break handler. Note the reset address is \$F[...]FFC0100. An exception will automatically switch the processor to the machine level operating mode. The break handler routine may redirect the exception to a lower level using the [REX](#) instruction.

The core maintains an internal eight level interrupt stack for each of the following:

Item Stacked	CSR reg	
instruction pointer	ip_stack	
operating level	ol_stack	available as a single CSR
privilege level	pl_stack	available as a single CSR
interrupt mask	im_stack	available as a single CSR

If further nesting of interrupts is required, the stacks may be copied to memory as they are available from CSR's.

On stack underflow a break exception is triggered.

Hardware interrupts will cause a BRK instruction to be inserted into the instruction stream.

Because instructions are fetched in bundles a hardware interrupt always intercepts at a bundle address, and always returns to a bundle address.

Instruction Format: BRK

H = 1 = software interrupt – return address is next instruction

H = 0 = hardware interrupt – return address is current instruction

IMask₄ = the priority level of the hardware interrupt, the priority level at time of interrupt is recorded in the instruction, the interrupt mask will be set to this level when the instruction commits. This field is not used for software interrupts and should be zero.

Cause Code = numeric code associated with the cause of the interrupt. The cause code is bitwise 'or'd with the value in register Rs1 to set the the cause CSR. Usually either one of the cause code field or Rs1 will be zero.

The empty instruction fields may be used to pass constant data to the break handler.

Operating Level	Address (If TVEC[0] contains \$FFFC0000)	
0	\$FFFC0000	Handler for operating level zero
1	\$FFFC0020	
2	\$FFFC0040	
3	\$FFFC0060	

CALL – Call Method

Description:

This instruction loads the instruction pointer with a constant value specified in the instruction. In addition, the address of the instruction following the CALL is stored in the return address register r61. This instruction may be used to implement subroutine calls. The target return address register is assumed to be r61 by the instruction. In order to use a different register, the JAL instruction must be used.

Instruction Format: CALL

This format has a 256GB range. The format modifies only instruction pointer bits 0 to 37. The high order IP bits are not affected. Note that with the use of a mmu this address range is often sufficient.

If an address range greater than 38 bits is required, then the JAL instruction must be used.

Execution Units: Branch Unit

Clock Cycles: 0.33

Notes:

Ideally the call instruction is placed in slot2 and targets an address in slot0. Many routines can be aligned on a bundle address. Since there are three instructions in a bundle the average extra space consumed by aligning routines on a bundle address is only 1.5 instruction words.

The call instruction executes immediately during the fetch stage of the core. This makes it much faster than a JAL.

CHK – Check Register Against Bounds

Description:

A register is compared to two values. If the register is outside of the bounds defined by Rs2 and Rs3 or an immediate value, then an exception will occur. Rs1 must be greater than or equal to Rs2 and Rs1 must be less than Rs3 or the immediate.

Instruction Format: CHK, CHKI

Clock Cycles: 0.33

Execution Units: Branch Unit

Exceptions: bounds check

Notes:

The system exception handler will typically transfer processing back to a local exception handler.

EXEC – Execute Code Buffer

Description:

Execute code from code buffer. The Codebuf₆ field (same field as Rs1) specifies the code buffer to use. Code buffers allow code to be adapted at run-time. This is useful as an alternative to self-modifying code when code must change at runtime. Some of the benefits of using a code buffer are the cache doesn't have to be invalidated, they are fast and easy to use.

Instruction Format: EXEC

Execution Units: Branch

Clock Cycles: Minimum 0.33 – depends on the instruction in the code buffer

Exceptions: As per instruction in code buffer.

Note:

The template to use for the instruction in the code buffer is found in the code buffer bits 40 to 46. The instruction to execute is found in bits 0 to 39. The template chosen should repeat the unit type of the instruction in the code buffer. For instance, if an integer alu instruction is being used the template should be the one for integer, integer, integer, or template 15h.

The EXEC instruction intercepts the instruction stream at the fetch phase and inserts the instruction from the code buffer rather than the cache output.

JAL – Jump-And-Link

Description:

Instruction Format:

This instruction loads the instruction pointer with the sum of a register and a constant value specified in the instruction. In addition, the address of the instruction following the JAL is stored in the specified target register. This instruction may be used to implement subroutine calls and returns.

Instruction Format: RI

Execution Units: Branch

Clock Cycles: 4

This instruction may not take effect until it is executed during the execute stage of the processor. It will cause the pipeline to flush which will dump following instructions from the queue. Depending on the size of the queue an average of ½ the number of entries in the queue will be dumped.

Exceptions: Branch target

JMP – Jump to Address

Description:

A jump is made to the address specified in the instruction.

Instruction Format: CALL

The format modifies only instruction pointer bits 0 to 37. The high order IP bits are not affected. This allows accessing code within a 256GB region of memory. Note that with the use of a mmu this address range is often sufficient.

Execution Units: Branch

Clock Cycles: 0.33

Exceptions: Branch target

Notes:

If an address range larger than 38 bits is required, then the value must be loaded into a register and the JAL instruction used.

The jump instruction executes immediately during the fetch stage of the core. This makes it much faster than a JAL.

JMP should be placed in slot2 for best performance.

NOP – No Operation

Description:

The NOP instruction doesn't perform any operation. NOP's are detected in the instruction fetch stage of the core and are not enqueued by the core. They do not occupy queue slots. Because NOPs don't occupy queue slots they may not be used to synchronize operations between instructions.

Instruction Format: NOP

Clock Cycles: 0.33

Execution Units: Branch

Exceptions: none

PFI – Poll for Interrupt

Description:

The poll for interrupt instruction polls the interrupt status lines and performs an interrupt service if an interrupt is present. Otherwise the PFI instruction is treated as a NOP operation. Polling for interrupts is performed by managed code. PFI provides a means to process interrupts at specific points in running software.

Note that the core records the address including the slot number of the PFI instruction as the exception address. This may be in the middle of an instruction bundle. In the case where an exception is present instructions in the bundle following the PFI instruction are not executed.

After completing the exception service routine, the core will return to the instruction following the PFI instruction, this may be in the middle of an instruction bundle.

Instruction Format: PFI

Clock Cycles: 0.33 (if no exception is present)

Execution Units: Branch

Exceptions: any outstanding hardware interrupts

RET – Return from Subroutine

Description:

This instruction performs a subroutine return by loading the instruction pointer with the contents of the return address register (r61). Additionally, the stack pointer is adjusted by a constant supplied in the instruction. The constant is zero extended to the left.

Instruction Format: RET

$PC = RA$

$SP = SP + \text{Immediate}$

Clock Cycles: 1 (more if predicted incorrectly).

Execution Units: Branch

Exceptions: none

Notes:

The registers used to implement the stack and return address are specified in the instruction format so different registers may be used. However, future versions of the core may not support this feature.

The RET instruction is detected and used at the fetch stage of the processor to update the RSB.

REX – Redirect Exception

Description:

This instruction redirects an exception from an operating level to a lower operating level and privilege level. If the target operating level is hypervisor then the hypervisor privilege level (1) is set. If the target operating level is supervisor, then one of the supervisor privilege levels must be chosen (2 to 6). This instruction if successful jumps to the target exception handler and does not return. If this instruction fails execution will continue with the next instruction.

This instruction may fail if exceptions are not enabled at the target level.

When redirecting the target privilege level is set to the bitwise 'or' of an immediate constant specified in the instruction and register Rs1. One of these two values should be zero. The result should be a value in the range 2 to 255. The instruction will not allow setting the privilege level numerically less than the operating level.

The location of the target exception handler is found in the trap vector register for that operating level (tvec[xx]).

The cause (cause) and bad address (badaddr) registers of the originating level are copied to the corresponding registers in the target level.

The REX instruction also specifies the interrupt mask level to set for further processing.

Attempting to redirect the operating level to the machine level (0) will be ignored. The instruction will be treated as a NOP with the exception of setting the interrupt mask register.

Instruction Format: REX

Tgt2	
0	not used
1	redirect to hypervisor level
2	redirect to supervisor level
3	not used

Clock Cycles: 4

Execution Units: Branch

Example:

```

REX 5,12,r0    ; redirect to supervisor handler, privilege level two
; If the redirection failed, exceptions were likely disabled at the target level.
; Continue processing so the target level may complete it's operation.
RTI            ; redirection failed (exceptions disabled ?)

```

Notes:

Since all exceptions are initially handled at the machine level the machine level handler must check for disabled lower level exceptions.

RTI – Return from Interrupt

Description:

Return from an interrupt or exception processing subroutine. The interrupted instruction pointer is loaded into the instruction pointer register. The internal interrupt stack is popped and the operating level, privilege level, interrupt mask level, and register set are reset to values before the exception occurred. Optionally a semaphore bit in the semaphore register is cleared. The least significant bit of the semaphore register (the reservation status bit) is always cleared by this instruction.

Instruction Format: RTI

$\text{Semaphore}[\text{Sema}_6[\text{Rs1}]] = 0$
 $\text{Semaphore}[0] = 0$

Clock Cycles: 8 minimum

Execution Units: Branch

SEI – Set Interrupt Mask

SEI #3

SEI \$v0,#7

Description:

The interrupt level mask is set to the value specified by the instruction. The value used is the bitwise or of the contents of register Rs1 and an immediate (M₄) supplied in the instruction. The assembler assumes a mask value of fifteen, masking all interrupts, if no mask value is specified. Usually either M₄ or Rs1 should be zero. The previous setting of the interrupt mask is stored in Rd.

Instruction Format: SEI

Operation:

Rd = im

im = M₄ | Rs1

SYNC -Synchronize

Description:

All instructions before the SYNC are completed and committed to the architectural state before instructions after the SYNC are issued. This instruction is used to ensure that the machine state is valid before subsequent instructions are executed.

Instruction Format: SYNC

Clock Cycles: 1 *varies depending on queue contents

Execution Units: Branch

Exceptions: none

Notes:

This instruction may be used with CSR register access as the core does not provide bypassing on the CSR registers. Issuing a sync instruction before reading a CSR will ensure that any outstanding updates to the CSR will be completed before the read.

Integer Unit

Integer Unit

This unit handles many different types of instructions associated with integer processing.

ABS – Absolute Value

Description:

This instruction takes the absolute value of a register and places the result in a target register.

Instruction Format: Integer R1

Clock Cycles: 0.33

Execution Units: Integer ALU

Operation:

```
If Rs1 < 0
    Rd = -Rs1
else
    Rd = Rs1
```

Exceptions: none

ADD - Addition

Description:

Add two values. The first operand must be in a register. The second operand may be in a register or may be an immediate value specified in the instruction.

Instruction Formats: Integer RI and R3

The R3 format adds three registers together.

Clock Cycles: 0.33

Execution Units: All ALU's

Exceptions: none

Notes:

ADDs1 – Addition Shifted 20 Bits

Description:

Add a register and an immediate value shifted to the left 20 bits. The immediate constant associated with the RI form of the instruction is sign extended to the left and zero extended to the right of the constant supplied in the instruction. This instruction allows building an 80-bit constant in a register when combined with other shifting instructions.

Instruction Format: Integer RI

Clock Cycles: 0.33

Execution Units: All Integer ALUs

Exceptions: none

ADDS2 – Addition Shifted 40 Bits

Description:

Add a register and an immediate value shifted to the left 40 bits. The immediate constant associated with the RI form of the instruction is sign extended to the left and zero extended to the right of the constant supplied in the instruction. This instruction allows building a 80-bit constant in a register when combined with other shifting instructions.

Instruction Format: Integer RI

Clock Cycles: 0.33

Execution Units: All Integer ALUs

Exceptions: none

ADDS3 – Addition Shifted 60 Bits

Description:

Add a register and an immediate value shifted to the left 60 bits. The immediate constant associated with the RI form of the instruction is zero extended to the right of the constant supplied in the instruction. This instruction allows building a 80-bit constant in a register when combined with other shifting instructions.

Instruction Format: Integer RI

Clock Cycles: 0.33

Execution Units: All Integer ALUs

Exceptions: none

AND – Bitwise And

Description:

Perform a bitwise ‘and’ operation between operands. For the immediate form, the immediate constant is one extended to the left before use.

Instruction Format: RI, R3

Clock Cycles: 0.33

Execution Units: Integer

Exceptions: none

ASL – Arithmetic Shift Left

Description:

Bits from the source register Rs1 are shifted left by the amount in register Rs2 or an immediate value. A zero is shifted into bit zero. A second operation which may be one of NOP, ADD or AND is performed between the first result and the value in register Rs3. The difference between this instruction and a SHL instruction is that ASL may cause an arithmetic overflow exception. SHL will never cause an exception.

The immediate constant for a shift is limited to six bits or the values 0 to 63. If it's required to shift by more than 63 bits two shift operations will be needed, or the value of the shift may come from a register.

Instruction Format: SHIFT, SHIFTI

Clock Cycles: 0.33

Execution Units: Integer

Exceptions:

An overflow exception may result if the bits shifted out from the MSB are not the same as the resulting sign bit and the exception is enabled in the AEC register. Exceptions are only caused by a word size operation.

Instruction Encoding Info:

Funct2 ₂	Operation
0	NOP
1	ADD
2	AND
3	reserved

ASR – Arithmetic Shift Right

Description:

Bits from the source register Rs1 are shifted right by the amount in register Rs2 or an immediate value. The sign bit is shifted into the most significant bits preserving the sign of the value. A second operation which may be one of NOP, ADD or AND is performed between the first result and the value in register Rs3.

The immediate constant for a shift is limited to six bits or the values 0 to 63. If it's required to shift by more than 63 bits two shift operations will be needed, or the value of the shift may come from a register.

Instruction Formats: SHIFT, SHIFTI

Clock Cycles: 0.33

Execution Units: Integer

Exceptions: none

Instruction Encoding Info:

Funct2 ₂	Operation
0	NOP
1	ADD
2	AND
3	reserved

AVG – Average Value

Description:

This instruction averages two signed values in registers Rs1 and Rs2 and places the result in a target Rd register. The average is calculated as the sum of Rs1 and Rs2 plus one for rounding, then arithmetically shift right by one.

Instruction Format: R1**Clock Cycles:** 0.33**Execution Units:** Integer**Operation:**

$$Rd = (Rs1 + Rs2 + 1) / 2$$

Exceptions: none

BFCHG – Bitfield Change

Description:

A bitfield in the source specified by Da is inverted, the result is copied to the target register. Bo specifies the bit offset. Bw specifies the bit width. The bit width and offset may either be contained in a register or an immediate value.

Instruction Format: BF

Rg ₃ Bit	
0	1 = Bo is a register spec, 0 = Bo is a six bit immediate
1	1 = Bw is a register spec, 0 = Bw is a six bit immediate
2	1

Clock Cycles: 0.33

Execution Units: Integer

Exceptions: none

BFCLR – Bitfield Clear

Description:

A bitfield is cleared in the target register. All bits are copied from a source Da (which is zero extended if an immediate value) except for bits identified by the bitfield which are set to zero in the target.

Instruction Format: BF

Rg ₃ Bit	
0	1 = Bo is a register spec, 0 = Bo is a six bit immediate
1	1 = Bw is a register spec, 0 = Bw is a six bit immediate
2	1 = Da is a register spec, 0 = Da is an sixteen bit immediate

Clock Cycles: 0.33

Execution Units: Integer

Exceptions: none

Notes:

Normally Da is a register which is the same as the target register Rt.

BFEXT – Bitfield Extract

Description:

A bitfield is extracted from the source register Da by shifting to the right and ‘and’ masking. The result is sign extended to the width of the machine. This instruction may be used to sign extend a value from an arbitrary bit position.

Instruction Format: BF

Rg ₃ Bit	
0	1 = Bo is a register spec, 0 = Bo is a six bit immediate
1	1 = Bw is a register spec, 0 = Bw is a six bit immediate
2	1 = Da is a register spec, 0 = Da is an sixteen bit immediate

Clock Cycles: 0.33

Execution Units: Integer

Exceptions: none

Notes:

While it is possible for Da to be a constant the instruction would not normally be used this way.

BFEXTU – Bitfield Extract

Description:

A bitfield is extracted from the source register Da by shifting to the right and ‘and’ masking. The result is zero extended to the width of the machine. This instruction may be used to zero extend a value from an arbitrary bit position.

Instruction Format: BF

Clock Cycles: 0.33

Execution Units: Integer

Exceptions: none

BFFFO – Bitfield Find First One

Description:

A bitfield contained in Da is searched beginning at the most significant bit to the least significant bit for a bit that is set. The index into the word of the bit that is set is stored in Rd. If no bits are set, then Rd is set equal to -1. To get the index into the bitfield of the set bit, subtract off the bitfield offset.

Instruction Format: BF

Rg ₃ Bit	
0	1 = Bo is a register spec, 0 = Bo is a six bit immediate
1	1 = Bw is a register spec, 0 = Bw is a six bit immediate
2	1 = Da is a register spec, 0 = Da is an sixteen bit immediate

Clock Cycles: 0.33

Execution Units: Integer

Exceptions: none

BFINSI – Bitfield Insert Immediate

Description:

A bitfield is inserted into the target register Rd by combining a value read from Rs3 with a constant shifted to the left. The bitfield may not be larger than six bits. To accommodate a larger field multiple instructions can be used, or a value loaded into a register and the BFINS instruction used.

Instruction Format: Integer BFI

Rg ₂ Bit	
0	1= Bo is a register spec, 0 = Bo is a six bit immediate
1	1 = Bw is a register spec, 0 = Bw is a six bit immediate

Rg_[1] bit should always be clear for this instruction

Clock Cycles: 1

Execution Units: ALU #0 Only

Exceptions: none

BMM – Bit Matrix Multiply

BMM Rd, Rs1, Rs2

Description:

The BMM instruction treats the bits of register Rs1 and Rs2 as an 8x8 bit matrix, performs a bit matrix multiply of the two registers and stores the result in the target register. An alternate mnemonic for this instruction is MOR. Only the least significant 64 bits of the registers are used.

Instruction Format: Integer R3

Func ₂	Function
0	MOR
1	MXOR
2	MORT (MOR transpose)
3	MXORT (MXOR transpose)

Operation:

for I = 0 to 7

for j = 0 to 7

$$Rt.bit[i][j] = (Ra[i][0] \& Rb[0][j]) \mid (Ra[i][1] \& Rb[1][j]) \mid \dots \mid (Ra[i][7] \& Rb[7][j])$$

Clock Cycles: 1

Execution Units: ALU #0 only

Exceptions: none

Notes:

The bits are numbered with bit 63 of a register representing $I,j = 0,0$ and bit 0 of the register representing $I,j = 7,7$.

BYTNDX – Byte Index

Description:

This instruction searches Rs1 for a byte specified by Rs2 or an immediate value and places the index of the byte into the target register Rd. If the byte is not found -1 is placed in the target register. A common use would be to search for a null byte. The index result may vary from -1 to +9. The index of the first found byte is returned (closest to zero).

Instruction Format: R3**Clock Cycles:** 0.33**Execution Units:** Integer**Operation:**
$$Rd = \text{Index of } (Rs2 \text{ in } Rs1)$$
Exceptions: none

CNTLO – Count Leading Ones

Description:

Count the number of leading ones (starting at the MSB) in Rs1 and place the count in the target register.

Instruction Format: R1**Clock Cycles:** 0.33**Execution Units:** Integer**Exceptions:** none

CNTLZ – Count Leading Zeros

Description:

Count the number of leading zeros (starting at the MSB) in Rs1 and place the count in the target register.

Instruction Format: R1**Clock Cycles:** 0.33**Execution Units:** Integer**Exceptions:** none

CNTPOP – Count Population

Description:

Count the number of one bits in Rs1 and place the count in the target register.

Instruction Format: R1**Clock Cycles:** 0.33**Execution Units:** Integer**Exceptions:** none

COM – One's Complement

Description:

This instruction takes the one's complement of a register and places the result in a target register. This is almost the same operation as exclusive or'ing with minus one, however the operation size may be set to operate on only a byte, wyde, tetra, penta, or octa value.

Instruction Format: Integer R1**Clock Cycles:** 0.33**Execution Units:** Integer ALU**Operation:**

$$Rd = \sim Rs1$$

Exceptions: none

CSR – Control and Status Access

Description:

The CSR instruction group provides access to control and status registers in the core. For the read-write operation the current value of the CSR is placed in the target register Rd then the CSR is updated from register Rs1. The CSR read / update operation is an atomic operation.

Instruction Format: CSR

Op ₂		Operation
0	CSRRD	Only read the CSR, no update takes place, Rs1 should be R0.
1	CSRRW	Both read and write the CSR
2	CSRRS	Read CSR then set CSR bits
3	CSRRC	Read CSR then clear CSR bits

CSRRS and CSRRC operations are only valid on registers that support the capability.

The OL₂ field is reserved to specify the operating level. Note that registers cannot be accessed by a lower operating level.

Regno ₁₂		Access	Description
001	HARTID	R	hardware thread identifier (core number)
002	TICK	R	tick count, counts every cycle from reset
030-037	TVEC	RW	trap vector handler address
048	EPC	RW	exceptioned pc, pc value at point of exception
044	STATUSL	RWSC	status register, contains interrupt mask, operating level
045	STATUSH	RW	status register bits 64 to 127
080-0BF	CODE	RW	code buffers
FF0	INFO	R	Manufacturer name
FF1	“	R	“
FF2	“	R	cpu class
FF3	“	R	“
FF4	“	R	cpu name
FF5	“	R	“
FF6	“	R	model number
FF7	“	R	serial number
FF8	“	R	cache sizes instruction (bits 40 to 79), data (bits 0 to 39)

Execution Units: Integer, the instruction may be available on only a single execution unit (not supported on all available integer units).

Clock Cycles: 0.33

Exceptions: privilege violation attempting to access registers outside of those allowed for the operating level.

DIF – Difference

Description:

This instruction computes the difference between two signed values in registers Rs1 and Rs2 and places the result in a target Rd register. The difference is calculated as the absolute value of Rs1 minus Rs2.

Instruction Format: R1**Clock Cycles:** 0.33**Execution Units:** Integer**Operation:**

$$Rd = \text{Abs}(Rs1 - Rs2)$$

Exceptions: none

DIV – Signed Division

Description:

Compute the quotient. The first operand must be in a register. The second operand may be in either a register or an immediate value specified in the instruction. The operands are treated as signed values and the result is a signed result.

Instruction Format: RI, R3

Clock Cycles: $84(n + 4)$ where n is the width

Execution Units: Integer

Exceptions: A divide by zero exception may occur if enabled in the AEC register.

DIVU – Unsigned Division

Description:

Compute the quotient value. The first operand must be in a register. The second operand may be in either a register or an immediate value specified in the instruction. The operands are treated as unsigned values and the result is an unsigned result.

Comment:

Unsigned division is often used in calculation of the difference between two pointer values.

Instruction Format: RI, R3

Clock Cycles: $84(n + 4)$ where n is the width

Execution Units: Integer

Exceptions: none

FXADD – Fixed Point Addition

Description:

Add two values assuming operands are signed fixed point numbers. The first operand must be in a register. The second operand may be in a register or may be an immediate value specified in the instruction. This instruction is an alternate mnemonic for the ADD instruction.

Instruction Format: RI, R3

Clock Cycles: 0.33

Execution Units: Integer

Exceptions: none

FXMUL – Fixed Point Multiplication

Description:

Multiply two values assuming operands are signed fixed point numbers. The first operand must be in a register. The second operand may be in a register or may be an immediate value specified in the instruction. The value is rounded towards plus infinity. The binary point is at the 40th bit position.

Instruction Format: RI, R3

Clock Cycles: 20

Execution Units: Integer

Exceptions: none

FXSUB – Fixed Point Subtraction

Description:

Subtract two values assuming operands are signed fixed point numbers. The first operand must be in a register. The second operand may be in a register or may be an immediate value specified in the instruction. This instruction is an alternate mnemonic for the SUB/ADD instruction.

Instruction Format: RI, R3

Clock Cycles: 0.33

Execution Units: Integer

Exceptions: none

ISPTR – Is Pointer

Description:

This instruction detects if the value in a register is a pointer and places the result in a target register. A pointer value is indicated if the top 24 bits of the value are equal to \$FFFF01.

Instruction Format:

Clock Cycles: 0.33

Execution Units: Integer

Operation:

```
If Ra[79,56] = $FFFF01
    Rt = 1
else
    Rt = 0
```

Exceptions: none

Notes:

MAJ – Majority Logic

Description:

Determines the majority logic bits of three values in registers Rs1, Rs2, and Rs3 and places the result in the target register Rd.

Instruction Format: R3

Clock Cycles: 0.33

Execution Units: Integer

Exceptions: none

Operation:

$$Rd = (Rs1 \& Rs2) | (Rs1 \& Rs3) | (Rs2 \& Rs3)$$

MAX – Maximum Value

Description:

Determines the maximum of three values in registers Rs1, Rs2, Rs3 and places the result in the target register Rd. The values are treated as signed values. In order to obtain the maximum value of two registers specify one of the registers twice.

MAX may be used to determine the lowest level privilege level. The lowest privilege level will have the highest value in registers.

Instruction Format: R3

Clock Cycles: 0.33

Execution Units: Integer

Operation:

```
IF Rs1 > Rs2 and Rs1 > Rs3
    Rd = Rs1
else if Rs2 > Rs3
    Rd = Rs2
else
    Rd = Rs3
```

Exceptions: none

MIN – Minimum Value

Description:

Determines the minimum of three values in registers Rs1, Rs2, Rs3 and places the result in the target register Rd. The values are treated as signed values. In order to obtain the minimum value of two registers specify one of the registers twice.

MIN may be used to determine the highest level privilege level. The highest privilege level will have the lowest value in registers.

Instruction Format: R3

Clock Cycles: 0.33

Execution Units: Integer

Operation:

```
IF Rs1 < Rs2 and Rs1 < Rs3
    Rd = Rs1
else if Rs2 < Rs3
    Rd = Rs2
else
    Rd = Rs3
```

Exceptions: none

MOD – Signed Modulus

Description:

Compute the modulus (remainder) value. The first operand must be in a register. The second operand may be in either a register or an immediate value specified in the instruction. The operands are treated as signed values and the result is a signed result.

Instruction Format: RI, R3**Clock Cycles:** $84(n + 4)$ where n is the width**Execution Units:** Integer**Exceptions:** A divide by zero exception may occur if enabled in the AEC register.

MODU – Unsigned Modulus

Description:

Compute the modulus (remainder) value. Both operands must be in registers. The operands are treated as unsigned values and the result is an unsigned result.

Instruction Format: R3**Clock Cycles:** $84(n + 4)$ where n is the width**Execution Units:** Integer**Exceptions:** none

MOV – Move register to register

Description:

This instruction moves one general purpose register to another.

Note that one does not normally want to move between integer and floating-point values directly. Instead usually a conversion is desired (ftoi - double to integer or itof - integer to double) for example.

Instruction Format: R1

Clock Cycles: 0.33

Execution Units: Integer

Exceptions: none

Notes:

MUL – Signed Multiply

Description:

Multiply two values. The first operand must be in a register Rs1. The second operand may be in a register Rs2 or may be an immediate value specified in the instruction. Both the operands are treated as signed values, the result is a signed result.

Perform a second function on the first result using Rs3. The function must be one of NOP, ADD or AND.

Instruction Format: RI, R3

Clock Cycles: 20

Execution Units: Integer

Exceptions: multiply overflow, if enabled

Instruction Encoding Info:

Funct2 ₂	Operation
0	NOP
1	ADD
2	AND
3	reserved

MULF – Fast Unsigned Multiply

Description:

Multiply two values. The first operand must be in a register. The second operand may be in a register or may be an immediate value specified in the instruction. Both the operands are treated as unsigned values. The result is an unsigned result. The fast multiply multiplies only the low order 24 bits of the first operand times the low order 16 bits of the second. The result is a 40-bit unsigned product. This instruction may be used to calculate array indices for small arrays.

Instruction Format: RI, R3

Clock Cycles: 0.33

Execution Units: Integer

Exceptions: none

MULH – Signed Multiply

Description:

Multiply two values. The first operand must be in a register. The second operand may be in a register or may be an immediate value specified in the instruction. Both the operands are treated as signed values, the result is a signed result. The high order bits of the product are returned.

Instruction Format: R3

Clock Cycles: 20

Execution Units: Integer

Exceptions: none

MULU – Unsigned Multiply

Description:

Multiply two values. The first operand must be in a register. The second operand may be in a register or may be an immediate value specified in the instruction. Both the operands are treated as unsigned values. The result is an unsigned result.

Comment:

Unsigned multiply is often used in address calculations for instance calculating array indexes.

Instruction Format: RI, R3

Clock Cycles: 20

Execution Units: Integer

Exceptions: none

MULUH – Unsigned Multiply

Description:

Multiply two values. The first operand must be in a register. The second operand may be in a register or may be an immediate value specified in the instruction. Both the operands are treated as unsigned values. The result is an unsigned result. The high order bits of the result are returned.

Instruction Format: R3**Clock Cycles:** 20**Execution Units:** Integer**Exceptions:** none

MUX – Multiplex

Description:

The MUX instruction performs a bit-by-bit copy of a bit of Rs2 to the target register Rd if the corresponding bit in Rs1 is set, or a copy of a bit from Rs3 if the corresponding bit in Ra is clear.

Instruction Format: R3**Clock Cycles:** 0.33**Execution Units:** Integer**Exceptions:** none

NAND – Bitwise Nand

Description:

Perform a bitwise and operation between two operands then invert the result. Both operands must be in registers.

Instruction Format: R3**Clock Cycles:** 0.33**Execution Units:** Integer**Exceptions:** none

NOR – Bitwise Nor

Description:

Perform a bitwise or operation between two operands then invert the result. Both operands must be in registers.

Instruction Format: R3**Clock Cycles:** 0.33**Execution Units:** Integer**Exceptions:** none

NOT – Logical Not

Description:

This instruction takes the logical ‘not’ value of a register and places the result in a target register. If the source register contains a non-zero value, then a zero is loaded into the target. Otherwise if the source register contains a zero a one is loaded into the target register.

Instruction Format: R1**Clock Cycles:** 0.33**Execution Units:** Integer**Operation:**

$$Rd = !Rs1$$

Exceptions: none**Notes:**

OR – Bitwise Or

Description:

Perform a bitwise or operation between operands. The immediate constant associated with the RI form of the instruction is zero extended to the size of the register.

Instruction Format: Integer RI and R3

Clock Cycles: 0.33

Execution Units: Integer

Exceptions: none

ORS1 – Bitwise Or Shifted 20 Bits

Description:

Perform a bitwise or operation between operands. The immediate constant associated with the RI form of the instruction is zero extended to the left and right of the constant supplied in the instruction. This instruction allows building a 80-bit constant in a register when combined with other 'or' and 'or' shifting instructions.

Instruction Format: Integer RI

Clock Cycles: 0.33

Execution Units: Integer

Exceptions: none

ORS2 – Bitwise Or Shifted 40 Bits

Description:

Perform a bitwise or operation between operands. The immediate constant associated with the RI form of the instruction is zero extended to the left and right of the constant supplied in the instruction. This instruction allows building a 80-bit constant in a register when combined with other 'or' and 'or' shifting instructions.

Instruction Format: Integer RI

Clock Cycles: 0.33

Execution Units: All Integer ALUs

Exceptions: none

ORS3 – Bitwise Or Shifted 60 Bits

Description:

Perform a bitwise or operation between operands. The immediate constant associated with the RI form of the instruction is zero extended to the right of the constant supplied in the instruction. This instruction allows building a 80-bit constant in a register when combined with other 'or' and 'or' shifting instructions.

Instruction Format: Integer RI

Clock Cycles: 0.33

Execution Units: All Integer ALUs

Exceptions: none

PTRDIF – Difference Between Pointers

Description:

Subtract two values then shift the result right. Both operands must be in a register. The top 20 bits of each value are masked off before the subtract. The right shift is provided to accommodate common object sizes. It may still be necessary to perform a divide operation after the PTRDIF to obtain an index into odd sized or large objects.

Instruction Format: R3

Operation:

$$Rt = (Ra - Rb) \gg Sc$$

Clock Cycles: 0.33

Execution Units: Integer

Exceptions:

None.

ROL – Rotate Left

Description:

Bits from the source register Ra are shifted left by the amount in register Rb or an immediate value. The most significant bit is shifted into bit zero.

For the sub-word forms the result is sign extended to 64 bits.

Instruction Format: SHIFT, SHIFTI**Clock Cycles:** 0.33**Execution Units:** Integer**Exceptions:** none

ROR – Rotate Right

Description:

Bits from the source register Rs1 are shifted right by the amount in register Rs2 or an immediate value. The bit zero is shifted into the most significant bits.

For the sub-word forms the result is sign extended to 64 bits.

Instruction Format: SHIFT, SHIFTI**Clock Cycles:** 0.33**Execution Units:** Integer**Exceptions:** none

SEQ – Set if Equal

Description:

The set instruction places a 1 or 0 in the target register based on the relationship between the two source operands. If operand Rs1 is equal to a second operand in a register (Rs2) or a constant supplied in the instruction, then the target register is set to a one, otherwise the target register is set to a zero. The operands may be either signed or unsigned values.

Instruction Format: RI, R3**Clock Cycles:** 0.33**Execution Units:** Integer**Exceptions:** none

SGE – Set if Greater Than or Equal

Description:

The set instruction places a 1 or 0 in the target register based on the relationship between the two source operands. If operand Rs1 is greater than or equal to a second operand in a register (Rs2) or a constant supplied in the instruction, then the target register is set to a one, otherwise the target register is set to a zero. The operands are treated as signed values.

Instruction Format: RI, R3

Clock Cycles: 0.33

Execution Units: Integer

Exceptions: none

SGEU – Set if Greater Than or Equal Unsigned

Description:

The set instruction places a 1 or 0 in the target register based on the relationship between the two source operands. If operand Rs1 is greater than or equal to a second operand in a register (Rs2) or a constant supplied in the instruction, then the target register is set to a one, otherwise the target register is set to a zero. The operands are treated as unsigned values.

Instruction Format: RI, R3

Clock Cycles: 0.33

Execution Units: Integer

Exceptions: none

SGT – Set if Greater Than

Description:

The set instruction places a 1 or 0 in the target register based on the relationship between the two source operands. If operand Rs1 is greater than a second operand in either a register (Rs2) or a constant supplied in the instruction, then the target register is set to a one, otherwise the target register is set to a zero. The operands are treated as signed values.

The register form of the instruction is just the register form of the SLT instruction with the operands switched.

Instruction Format: RI, R3

Clock Cycles: 0.33

Execution Units: Integer

Exceptions: none

SGTU – Set if Greater Than Unsigned

Description:

The set instruction places a 1 or 0 in the target register based on the relationship between the two source operands. If operand Rs1 is greater than a second operand in either a register (Rs2) or a constant supplied in the instruction, then the target register is set to a one, otherwise the target register is set to a zero. The operands are treated as unsigned values.

The register form of the instruction is just the register form of the SLTU instruction with the operands switched.

Instruction Format: RI, R3

Clock Cycles: 0.33

Execution Units: Integer

Exceptions: none

SHL – Shift Left

Description:

Bits from the source register Rs1 are shifted left by the amount in register Rs2 or an immediate value. Zeros are shifted into the least significant bits. A second operation which may be one of NOP, ADD or AND is performed between the first result and the value in register Rs3.

The immediate constant for a shift is limited to six bits or the values 0 to 63. If it's required to shift by more than 63 bits two shift operations will be needed, or the value of the shift may come from a register.

Instruction Format: SHIFT, SHIFTI

Clock Cycles: 0.33

Execution Units: Integer

Exceptions: none

Instruction Encoding Info:

Funct2 ₂	Operation
0	NOP
1	ADD
2	AND
3	reserved

SLE – Set if Less Than or Equal

Description:

The set instruction places a 1 or 0 in the target register based on the relationship between the two source operands. If operand Rs1 is less than or equal to a second operand in a register (Rs2) or a constant supplied in the instruction, then the target register is set to a one, otherwise the target register is set to a zero. The operands are treated as signed values.

The instruction may also be used to test for greater than or equal by swapping the operands around.

Instruction Format: RI, R3

Clock Cycles: 0.33

Execution Units: Integer

Exceptions: none

SLEU – Set if Less Than or Equal Unsigned

Description:

The set instruction places a 1 or 0 in the target register based on the relationship between the two source operands. If operand Rs1 is less than or equal to a second operand in a register (Rs2) or a constant supplied in the instruction, then the target register is set to a one, otherwise the target register is set to a zero. The operands are treated as unsigned values.

The instruction may also be used to test for greater than or equal by swapping the operands around.

Instruction Format: RI, R3

Clock Cycles: 0.33

Execution Units: Integer

Exceptions: none

SLT – Set if Less Than

Description:

The set instruction places a 1 or 0 in the target register based on the relationship between the two source operands. If operand Rs1 is less than a second operand in either a register (Rs2) or a constant supplied in the instruction, then the target register is set to a one, otherwise the target register is set to a zero. The operands are treated as signed values.

Instruction Format: RI, R3

Clock Cycles: 0.33

Execution Units: Integer

Exceptions: none

SLTU – Set if Less Than

Description:

The set instruction places a 1 or 0 in the target register based on the relationship between the two source operands. If operand Rs1 is less than a second operand in either a register (Rs2) or a constant supplied in the instruction, then the target register is set to a one, otherwise the target register is set to a zero. The operands are treated as unsigned values.

Instruction Format: RI, R3

Clock Cycles: 0.33

Execution Units: Integer

Exceptions: none

SNE – Set if Not Equal

Description:

The set instruction places a 1 or 0 in the target register based on the relationship between the two source operands. If operand Rs1 is not equal to a second operand in a register (Rs2) or a constant supplied in the instruction, then the target register is set to a one, otherwise the target register is set to a zero. The operands may be either signed or unsigned values.

Instruction Format: RI, R3

Clock Cycles: 0.33

Execution Units: Integer

Exceptions: none

SXB – Sign Extend Byte

Description:

Sign extend a byte (8 bits) as a signed value to the full width of the machine.

Instruction Format: R1

Clock Cycles: 0.33

Execution Units: Integer

Exceptions: none

SXO – Sign Extend Octa Byte

Description:

Sign extend an octa byte (64 bits) as a signed value to the full width of the machine.

Instruction Format: R1

Clock Cycles: 0.33

Execution Units: Integer

Exceptions: none

SXP – Sign Extend Penta Byte

Description:

Sign extend a penta byte (40 bits) as a signed value to the full width of the machine.

Instruction Format: R1

Clock Cycles: 0.33

Execution Units: Integer

Exceptions: none

SXT – Sign Extend Tetra Byte

Description:

Sign extend a tetra byte (32 bits) as a signed value to the full width of the machine.

Instruction Format: R1

Clock Cycles: 0.33

Execution Units: Integer

Exceptions: none

SXW – Sign Extend Wyde

Description:

Sign extend a wyde (16 bits) as a signed value to the full width of the machine.

Instruction Format: R1

Clock Cycles: 0.33

Execution Units: Integer

Exceptions: none

WYDNDX – Wyde Index

Description:

This instruction searches Rs1, which is treated as an array of five wydes, for a wyde value specified by Rs2 or an immediate value and places the index of the wyde into the target register Rd. If the wyde is not found -1 is placed in the target register. A common use would be to search for a null wyde. The index result may vary from -1 to +4. The index of the first found wyde is returned (closest to zero).

Instruction Format: R3**Clock Cycles:** 0.33**Execution Units:** Integer**Operation:**
$$Rd = \text{Index of } (Rs2 \text{ in } Rs1)$$
Exceptions: none

XNOR – Bitwise Exclusive Nor

Description:

Perform a bitwise exclusive or operation between two operands then invert the result. Both operands must be in registers.

Instruction Format: R3**Clock Cycles:** 0.33**Execution Units:** Integer**Exceptions:** none

XOR – Bitwise Exclusive Or

Description:

Perform a bitwise exclusive or operation between operands. The first operand must be in a register. The second operand may be a register or immediate value.

Instruction Format: RI, R3**Clock Cycles:** 0.33**Execution Units:** Integer**Exceptions:** none

ZXB – Zero Extend Byte

Description:

Zero extend a byte (8 bits) as an unsigned value to the full width of the machine.

Instruction Format: R1

Clock Cycles: 0.33

Execution Units: Integer

Exceptions: none

ZXO – Zero Extend Octa Byte

Description:

Zero extend an octa byte (64 bits) as an unsigned value to the full width of the machine.

Instruction Format: R1

Clock Cycles: 0.33

Execution Units: Integer

Exceptions: none

ZXP – Zero Extend Penta Byte

Description:

Zero extend a penta byte (40 bits) as an unsigned value to the full width of the machine.

Instruction Format: R1

Clock Cycles: 0.33

Execution Units: Integer

Exceptions: none

ZXT – Zero Extend Tetra

Description:

Zero extend a tetra byte (32 bits) as an unsigned value to the full width of the machine.

Instruction Format: R1

Clock Cycles: 0.33

Execution Units: Integer

Exceptions: none

ZXW – Zero Extend Wyde

Description:

Zero extend a wyde byte (16 bits) as an unsigned value to the full width of the machine.

Instruction Format: R1

Clock Cycles: 0.33

Execution Units: Integer

Exceptions: none

Memory Unit

Preload Instructions

Issuing a load instruction with a target register of r0 causes the cache line to be loaded without updating the register file. If an exception occurs during the load, it will be ignored.

AMO – Atomic Memory Operation

Description:

The atomic memory operations read from memory addressed by the Rs1 register and store the value in Rd. As a second step the value from memory is combined with the value in register Rs2 or an immediate constant according to one of the available functions then stored back into the memory addressed by Rs1.

Instruction Format: AMO, AMOI

Funct ₄	Mnemonic	Operation Performed	
00	amoswap	swap	memory[Ra] = Rb
01	amoadd	addition	memory[Ra] = memory[Ra] + Rb
02	amoand	bitwise and	memory[Ra] = memory[Ra] & Rb
03	amoor	bitwise or	memory[Ra] = memory[Ra] Rb
04	amoxor	bitwise exclusive or	memory[Ra] = memory[Ra] ^ Rb
05	amoshl	shift left	memory[Ra] = memory[Ra] << Rb
06	amoshr	shift right	memory[Ra] = memory[Ra] >> Rb
07	amomin	minimum	memory[Ra] = memory[Ra] < Rb ? memory[Ra] : Rb
08	amomax	maximum	memory[Ra] = memory[Ra] > Rb ? memory[Ra] : Rb
09	amominu	minimum unsigned	memory[Ra] = memory[Ra] < Rb ? memory[Ra] : Rb
0A	amomaxu	maximum unsigned	memory[Ra] = memory[Ra] > Rb ? memory[Ra] : Rb

Sz ₃	
0	Byte
1	Wbyte
2	Tetra
3	Penta
4	Octa
5	Deci
6	reserved
7	reserved

Acquire and release bits determine the ordering of memory operations.

A = acquire – 1 = no following memory operations can take place before this one

R = release – 1 = this memory operation cannot take place before prior ones.

All combinations of A, R are allowed.

CACHE – Cache Command

CACHE Cmd, d[Rn]

CACHE Cmd, d[Ra + Rc * scale]

Description:

This instruction commands the cache controller to perform an operation. Commands are summarized in the command table below. Commands may be issued to both the instruction and data cache at the same time.

Instruction Formats: CACHE

Commands:

ICCmd ₂	Mne.	Operation
0	NOP	no operation
1	invline	invalidate line associated with given address
2	invall	invalidate the entire cache (address is ignored)

DCCmd ₃	Mne.	Operation
0	NOP	no operation
1	enable	enable cache (instruction cache is always enabled)
2	disable	not valid for the instruction cache
3	invline	invalidate line associated with given address
4	invall	invalidate the entire cache (address is ignored)

Operation:

Register Indirect with Displacement Form

$$\text{Line} = \text{round}_{32}(\text{sign extend}(\text{memory}[\text{displacement} + \text{Ra}]))$$

Register-Register Form

$$\text{Line} = \text{round}_{32}(\text{sign extend}(\text{memory}[\text{Ra} + \text{Rc} * \text{scale}]))$$

Notes:

Sc ₂ Code	Multiply By
0	1
1	2
2	4
3	8

LDB – Load Byte

Description:

This instruction loads a byte (8 bit) value from memory. The value is sign extended to 80 bits when placed in the target register.

Instruction Formats: ML, MLX

Clock Cycles: 4 minimum depending on memory access time

Sc ₃	Scale Rb By
0	1
1	2
2	4
3	8
4	16
5	5
6	10
7	15

LDBU – Load Unsigned Byte

Description:

This instruction loads a byte (8 bit) value from memory. The value is zero extended to 80 bits when placed in the target register.

Instruction Formats: ML, MLX

Clock Cycles: 4 minimum depending on memory access time

Sc ₃	Scale Rb By
0	1
1	2
2	4
3	8
4	16
5	5
6	10
7	15

LDD – Load Deci Byte (80 bits)

Description:

This instruction loads a byte (80 bit) value from memory

Instruction Formats: ML, MLX

Clock Cycles: 4 minimum depending on memory access time

Sc ₃	Scale Rb By
0	1
1	2
2	4
3	8
4	16
5	5
6	10
7	15

LDDR – Load Deci Byte and Reserve

Description:

This instruction loads a byte (80 bit) value from memory and places a reservation on the address range associated with the given address. The region of memory which is reserved depends on what is supported by the external memory system hardware. Typically a 128 bit region or larger would be reserved.

Instruction Formats: ML, MLX

Clock Cycles: 4 minimum depending on memory access time

Sc ₃	Scale Rb By
0	1
1	2
2	4
3	8
4	16
5	5
6	10
7	15

LDFD – Load Float Deci-Byte (80 bits)

Description:

This instruction loads a deci-byte (80 bit) value from memory into a floating-point register.

Instruction Formats: ML, MLX

Clock Cycles: 4 minimum depending on memory access time

Sc ₃	Scale Rb By
0	1
1	2
2	4
3	8
4	16
5	5
6	10
7	15

LDO – Load Octa Byte (64 bits)

Description:

This instruction loads a byte (64 bit) value from memory. The value is sign extended to 80 bits when placed in the target register.

Instruction Formats: ML, MLX

Clock Cycles: 4 minimum depending on memory access time

Sc ₃	Scale Rb By
0	1
1	2
2	4
3	8
4	16
5	5
6	10
7	15

LDOU – Load Unsigned Octa Byte

Description:

This instruction loads a byte (64 bit) value from memory. The value is zero extended to 80 bits when placed in the target register.

Instruction Formats: ML, MLX

Clock Cycles: 4 minimum depending on memory access time

Sc ₃	Scale Rb By
0	1
1	2
2	4
3	8
4	16
5	5
6	10
7	15

LDP – Load Penta Byte (40 bits)

Description:

This instruction loads a penta-byte (40 bit) value from memory. The value is sign extended to 80 bits when placed in the target register.

Instruction Formats: ML, MLX

Clock Cycles: 4 minimum depending on memory access time

Sc ₃	Scale Rb By
0	1
1	2
2	4
3	8
4	16
5	5
6	10
7	15

LDPU – Load Unsigned Penta Byte

Description:

This instruction loads a penta-byte (40 bit) value from memory. The value is zero extended to 80 bits when placed in the target register.

Instruction Formats: ML, MLX

Clock Cycles: 4 minimum depending on memory access time

Sc ₃	Scale Rb By
0	1
1	2
2	4
3	8
4	16
5	5
6	10
7	15

LDT – Load Tetra Byte (32 bits)

Description:

This instruction loads a penta-byte (32 bit) value from memory. The value is sign extended to 80 bits when placed in the target register.

Instruction Formats: ML, MLX

Clock Cycles: 4 minimum depending on memory access time

Sc ₃	Scale Rb By
0	1
1	2
2	4
3	8
4	16
5	5
6	10
7	15

LDTU – Load Unsigned Tetra Byte

Description:

This instruction loads a penta-byte (32 bit) value from memory. The value is zero extended to 80 bits when placed in the target register.

Instruction Formats: ML, MLX

Clock Cycles: 4 minimum depending on memory access time

Sc ₃	Scale Rb By
0	1
1	2
2	4
3	8
4	16
5	5
6	10
7	15

LDW – Load Wyde (16 bits)

Description:

This instruction loads a wyde (16 bit) value from memory. The value is sign extended to 80 bits when placed in the target register.

Instruction Formats: ML, MLX

Clock Cycles: 4 minimum depending on memory access time

Sc ₃	Scale Rb By
0	1
1	2
2	4
3	8
4	16
5	5
6	10
7	15

LDWU – Load Unsigned Wyde

Description:

This instruction loads a byte (16 bit) value from memory. The value is zero extended to 80 bits when placed in the target register.

Instruction Formats: ML, MLX

Clock Cycles: 4 minimum depending on memory access time

Sc ₃	Scale Rb By
0	1
1	2
2	4
3	8
4	16
5	5
6	10
7	15

LEA – Load Effective Address

Description:

This instruction loads an address value into a register. The address value is calculated in the same manner as a load / store instruction. The upper 24 bits of the register are automatically set to \$FFFF01 to indicate a pointer value is in the register.

Instruction Format: ML, MLX

Clock Cycles: 0.33

Sc ₂	Scale Rb By
0	1
1	2
2	4
3	8

Execution Units: Memory Load

Operation:

Exceptions: none

Notes:

LINK – Link Stack

Description:

This instruction pushes the frame pointer (r62) to the stack then increments the stack pointer and transfers the incremented stack pointer to the frame pointer register. Finally, an amount specified in the instruction is subtracted from the stack pointer.

Instruction Format: LINK

Operation:

$$\text{Memory}_{10}[\text{SP} - 10] = \text{FP}$$
$$\text{SP} = \text{SP} - \text{amt} - 10$$
$$\text{FP} = \text{SP}$$

Clock Cycles: 4 minimum, depending on memory access time

Execution Units: Memory

Notes:

The instruction explicitly encodes the stack pointer register r63, if desired a different register may be chosen. However, stack bounds checking is available with register r63.

MEMDB –Memory Data Barrier

Description:

All memory instructions before the MEMDB are completed and committed to the architectural state before memory instructions after the MEMDB are issued. This instruction is used to ensure that the memory state is valid before subsequent instructions are executed.

Instruction Format: MEMDB

Clock Cycles: varies depending on queue contents

Execution Units: Memory Store

Operation:

Exceptions: none

MEMSB –Memory Synchronization Barrier

Description:

This instruction is similar to the SYNC instruction except that it applies only to memory operations. All instructions before the MEMSB are completed and committed to the architectural state before memory instructions after the MEMSB are issued. This instruction is used to ensure that the memory state is valid before subsequent instructions are executed.

Instruction Format: MEMSB

Clock Cycles: varies depending on queue contents

Execution Units: Memory Store

Operation:

Exceptions: none

POP – Pop (80 bits)

Description:

This instruction moves a deci-byte value from the stack to a register then increments the stack pointer.

Instruction Format: POP

Operation:

$$\begin{aligned} \text{Rd} &= \text{Memory}_{10}[\text{SP}] \\ \text{SP} &= \text{SP} + 10 \end{aligned}$$

Clock Cycles: 4 minimum, depending on memory access time

Execution Units: Memory

Notes:

The instruction explicitly encodes the stack pointer register r63, if desired a different register may be chosen. However, stack bounds checking is available with register r63.

PUSH – Push (80 bits)

Description:

This instruction decrements the stack pointer and stores a deci byte (80 bit) value to stack memory. The value pushed onto the stack may come from either a register or a constant value defined in the instruction.

Instruction Format: PUSH, PUSHC

Operation:

$$\text{Memory}_8[\text{SP} - 10] = \text{Rs2}$$
$$\text{SP} = \text{SP} - 10$$

Clock Cycles: 4 minimum, depending on memory access time

Execution Units: Memory Store

Notes:

Stores always write through to memory and therefore take a significant number of clock cycles before they are ready to be committed. Exceptions are checked for during the execution of a store operation.

The instruction explicitly encodes the stack pointer register r63, if desired a different register may be chosen. However, stack bounds checking is available with register r63.

STB – Store Byte (8 bits)

Description:

This instruction stores a byte (8 bit) value to memory.

Instruction Format: MS, MSX

Operation:

$$\text{Memory}_8[\text{Rs1} + \text{immediate}] = \text{Rs2}$$

Clock Cycles: 4 minimum depending on memory access time

The number of clock cycles required by this instruction may be “hidden” because the core will continue to process instructions while the store is being processed.

Sc ₃	Scale Rb By
0	1
1	2
2	4
3	8
4	16
5	5
6	10
7	15

Execution Units: Memory Store

Exceptions: DBE, TLB, WRV

Notes:

A store instruction will only be issued by the core if it is known that it will be able to commit without being interrupted by a change of program flow due to a previous instruction. Stores always write through to memory and therefore take a significant number of clock cycles before they are ready to be committed. Exceptions are checked for during the execution of a store operation.

STD – Store Deci (80 bits)

Description:

This instruction stores a deci byte (80 bit) value to memory.

Instruction Format: MS, MSX

Operation:

$\text{Memory}_{80}[\text{Rs1} + \text{immediate}] = \text{Rs2}$

Clock Cycles: 4 minimum depending on memory access time

The number of clock cycles required by this instruction may be “hidden” because the core will continue to process instructions while the store is being processed.

Sc ₃	Scale Rb By
0	1
1	2
2	4
3	8
4	16
5	5
6	10
7	15

Execution Units: Memory Store

Notes:

A store instruction will only be issued by the core if it is known that it will be able to commit without being interrupted by a change of program flow due to a previous instruction. Stores always write through to memory and therefore take a significant number of clock cycles before they are ready to be committed. Exceptions are checked for during the execution of a store operation.

STFD – Store Float Deci-byte (80 bits)

Description:

This instruction stores an deci byte (80 bit) value to memory.

Instruction Format: MS, MSX

Operation:

$$\text{Memory}_{80}[\text{Rs1} + \text{immediate}] = \text{Rs2}$$

Clock Cycles: 4 minimum depending on memory access time

The number of clock cycles required by this instruction may be “hidden” because the core will continue to process instructions while the store is being processed.

Sc ₃	Scale Rb By
0	1
1	2
2	4
3	8
4	16
5	5
6	10
7	15

Execution Units: Memory Store

Notes:

A store instruction will only be issued by the core if it is known that it will be able to commit without being interrupted by a change of program flow due to a previous instruction. Stores always write through to memory and therefore take a significant number of clock cycles before they are ready to be committed. Exceptions are checked for during the execution of a store operation.

STO – Store Octa (64 bits)

Description:

This instruction stores an octa byte (64 bit) value to memory.

Instruction Format: MS, MSX

Operation:

$$\text{Memory}_{64}[\text{Rs1} + \text{immediate}] = \text{Rs2}$$

Clock Cycles: 4 minimum depending on memory access time

The number of clock cycles required by this instruction may be “hidden” because the core will continue to process instructions while the store is being processed.

Sc ₃	Scale Rb By
0	1
1	2
2	4
3	8
4	16
5	5
6	10
7	15

Execution Units: Memory Store

Notes:

A store instruction will only be issued by the core if it is known that it will be able to commit without being interrupted by a change of program flow due to a previous instruction. Stores always write through to memory and therefore take a significant number of clock cycles before they are ready to be committed. Exceptions are checked for during the execution of a store operation.

STP – Store Penta (40 bits)

Description:

This instruction stores a penta (40 bit) value to memory. This operation is $\frac{1}{2}$ the 80-bit word size of the machine, and also the size of an instruction in an instruction bundle.

Instruction Format: MS, MSX

Operation:

$\text{Memory}_{40}[\text{Rs1} + \text{immediate}] = \text{Rs2}$

Clock Cycles: 4 minimum depending on memory access time

The number of clock cycles required by this instruction may be “hidden” because the core will continue to process instructions while the store is being processed.

Sc ₃	Scale Rb By
0	1
1	2
2	4
3	8
4	16
5	5
6	10
7	15

Execution Units: Memory Store

Notes:

A store instruction will only be issued by the core if it is known that it will be able to commit without being interrupted by a change of program flow due to a previous instruction. Stores always write through to memory and therefore take a significant number of clock cycles before they are ready to be committed. Exceptions are checked for during the execution of a store operation.

STT – Store Tetra (32 bits)

Description:

This instruction stores a tetra (32 bit) value to memory.

Instruction Format: MS, MSX

Operation:

$$\text{Memory}_{32}[\text{Rs1} + \text{immediate}] = \text{Rs2}$$

Clock Cycles: 4 minimum depending on memory access time

The number of clock cycles required by this instruction may be “hidden” because the core will continue to process instructions while the store is being processed.

Sc ₃	Scale Rb By
0	1
1	2
2	4
3	8
4	16
5	5
6	10
7	15

Execution Units: Memory Store

Notes:

A store instruction will only be issued by the core if it is known that it will be able to commit without being interrupted by a change of program flow due to a previous instruction. Stores always write through to memory and therefore take a significant number of clock cycles before they are ready to be committed. Exceptions are checked for during the execution of a store operation.

STW – Store Wyde (16 bits)

Description:

This instruction stores a wyde (16 bit) value to memory.

Instruction Format: MS, MSX

Operation:

$$\text{Memory}_{16}[\text{Rs1} + \text{immediate}] = \text{Rs2}$$

Clock Cycles: 4 minimum depending on memory access time

The number of clock cycles required by this instruction may be “hidden” because the core will continue to process instructions while the store is being processed.

Sc ₃	Scale Rb By
0	1
1	2
2	4
3	8
4	16
5	5
6	10
7	15

Execution Units: Memory Store

Notes:

A store instruction will only be issued by the core if it is known that it will be able to commit without being interrupted by a change of program flow due to a previous instruction. Stores always write through to memory and therefore take a significant number of clock cycles before they are ready to be committed. Exceptions are checked for during the execution of a store operation.

TLB – TLB Command

Description:

The command is executed on the TLB unit. The command results are placed in internal TLB registers which can be read or written using TLB command instruction. If the operation is a read register operation, then the register value is placed into Rd. If the operation is a write register operation, then the value for the register comes from Rs1. Otherwise the Rs1/Rd field in the instruction is ignored.

This instruction is only available at the machine operating level.

Instruction Format: TLB

Clock Cycles: 3

It typically takes several cycles for the TLB to process the command.

Tn₄ – This field identifies which TLB register is being read or written.

Tn ₄		Assembler
0	Wired	Wired
1	Index	Index
2	Random	Random
3	Page Size	PageSize
4	Virtual page	VirtPage
5	Physical page	PhysPage
7	ASID	ASID
8	Miss address	MA
9	reserved	
10	Page Table Address	PTA
11	Page Table Control	PTC
12	Aging frequency control	AFC

TLB Commands

Cmd ₅	Description	Assembler
0	No operation	
1	Probe TLB entry	TLBPB
2	Read TLB entry	TLBRD
3	Write TLB entry corresponding to random register	TLBWR
4	Write TLB entry corresponding to index register	TLBWI
5	Enable TLB	TLBEN
6	Disable TLB	TLBDIS
7	Read register	TLBRDREG
8	Write register	TLBWRREG
9	Invalidate all entries	TLBINV
10	Get age count	TLBRDAGE
11	Set age count	TLBWRAGE

Probe TLB – The TLB will be tested to see if an address translation is present.

Read TLB – The TLB entry specified in the index register will be copied to TLB holding registers.

Write Random TLB – A random TLB entry will be written into from the TLB holding registers.

Write Indexed TLB – The TLB entry specified by the index register will be written from the TLB holding registers.

Disable TLB – TLB address translation is disabled so that the physical address will match the supplied virtual address.

Enable TLB – TLB address translation is enabled. Virtual address will be translated to physical addresses using the TLB lookup tables.

The TLB will automatically update the miss address registers when a TLB miss occurs only if the registers are zero to begin with. System software must reset the registers to zero after a miss is processed. This mechanism ensures the first miss that occurs is the one that is recorded by the TLB.

PageTableAddr – This is a scratchpad register available for use to store the address of the page table.

PageTableCtrl – This is a scratchpad register available for use to store control information associated with the page table.

Notes:

The TLB automatically tracks reference counts and ages address mappings.

UNLK – Unlink Stack

Description:

This instruction is used to unlink the stack frame by restoring the stack pointer from the frame pointer, then popping the frame pointer off the stack.

Instruction Format: UNLK

Operation:

$$\begin{aligned} SP &= FP \\ FP &= \text{Memory}_{10}[SP] \\ SP &= SP + 10 \end{aligned}$$

Clock Cycles: 4 minimum, depending on memory access time

Execution Units: Memory

Notes:

The instruction explicitly encodes the stack pointer register r63, if desired a different register may be chosen. However, stack bounds checking is available with register r63.

Floating Point Instruction Description

Overview

The floating-point unit provides basic floating-point operations including addition, subtraction, multiplication, division, square root, and float to integer and integer to float conversions. The core contains two identical floating-point units. Only extended double precision floating-point operations are supported.

Unlike some FP implementations, the core allows the use of immediate constants supplied in the instruction stream as an operand. The FLTI format instructions use a large immediate constant as the second operand to the instruction.

The precision field (prec_4) should be set to 2.

The rounding mode is normally specified directly in the instruction. However if the instruction indicates to use dynamic rounding mode then the rounding mode in the floating-point control and status register is used.

Internally the core maintains four extra mantissa bits, these extra bits are truncated on stores and zeroed on loads. However, they are used during calculations, as long as results remain in registers the four extra bits will provide four more bits of precision.

Representation

The floating-point format is similar to an IEEE-754 representation for extended double precision. Briefly,

Extended Double Precision Format:

79	78	77	64	63	0
S_M	S_E	Exponent	Mantissa		

S_M – sign of mantissa

S_E – sign of exponent

The exponent and mantissa are both represented as two's complement numbers, however the sign bit of the exponent is inverted.

S_e EEEEEEEEEE	
1111111111	Maximum exponent
....	
0111111111	exponent of zero
....	
0000000000	Minimum exponent

The exponent ranges from -16384 to +16383 for extended double precision numbers

Prec_4	Precision
0	Half (20 bit)
1	Single (40 bit)
2	Double (80 bit)

3	reserved
4	reserved
5 to 7	reserved

F32ToF80 – Convert 32-bit to 80-bit

Description:

Convert the single precision value (32 bits) in Rs1 into a floating-point double value (80 bits) and place the result into target register Rd.

Instruction Format: FLT1

Clock Cycles: 0.33

Execution Units: Floating Point

F40ToF80 / FCVTSD – Convert Single to Double

Description:

Convert the single precision value (40 bits) in Rs1 into a floating-point double value (80 bits) and place the result into target register Rd.

Instruction Format: FLT1

Clock Cycles: 3

Execution Units: Floating Point

F64ToF80 – Convert 64-bit to 80-bit

Description:

Convert the double precision value (64 bits) in Rs1 into a floating-point double value (80 bits) and place the result into target register Rd.

Instruction Format: FLT1

Clock Cycles: 0.33

Execution Units: Floating Point

F80ToF32 – Convert 80-bit to 32-bit

Description:

Convert the double precision value (80 bits) in Rs1 into a floating-point single value (32 bits) and place the result into target register Rd. Not all values are possible to convert to 32-bits, some precision is lost. Values outside the range of a 32-bit float may be set to either infinity or zero. The 64-bit mantissa is truncated to 23 bits. Nans are propagated.

Instruction Format: FLT1**Clock Cycles:** 0.33**Execution Units:** Floating Point

F80ToF40 – Convert 80-bit to 40-bit

Description:

Convert the double precision value (80 bits) in Rs1 into a floating-point single value (40 bits) and place the result into target register Rd. Not all values are possible to convert to 40-bits, some precision is lost. Values outside the range of a 40-bit float may be set to either infinity or zero. The 40-bit mantissa is truncated to 29 bits. Nans are propagated.

Instruction Format: FLT1**Clock Cycles:** 0.33**Execution Units:** Floating Point

F80ToF64 – Convert 80-bit to 64-bit

Description:

Convert the double precision value (80 bits) in Rs1 into a floating-point double value (64 bits) and place the result into target register Rd. Not all values are possible to convert to 64-bits, some precision is lost. Values outside the range of a 64-bit float may be set to either infinity or zero. The 64-bit mantissa is truncated to 52 bits. Nans are propagated.

Instruction Format: FLT1

Clock Cycles: 0.33

Execution Units: Floating Point

FABS – Floating Absolute Value

Description:

This instruction is an alternate mnemonic for the FSGNCPY instruction where Rs2 is zero. It effectively takes the absolute value of the floating-point number in register Rs1 and places the result into target register Rd. The sign bit (bit 79) of the register is set to zero. No rounding of the number occurs.

Instruction Format: FLT2

Clock Cycles: 0.33

Execution Units: Floating Point

FADD – Floating point addition

Description:

Add two floating point numbers in registers Rs1 and Rs2 or an immediate constant and place the result into target register Rd.

Instruction Format: FLT2, FLTI

Clock Cycles: 6

Execution Units: Floating Point

FASGE - Float Set if Absolute Value Greater Than or Equal

Description:

The register compare instruction compares the absolute values in two registers or a register and an immediate constant as floating-point values for greater than or equal and sets the target register as a result.

Instruction Format: FLT2, FLTI

Clock Cycles: 0.33

Execution Units: Floating Point

Operation:

```
if abs(Rs1) >= abs(Rs2)
    Rd = true
else
    Rd = false
```

FBcc – Conditional Branch

Description:

If the branch condition is true, a target address is loaded into the instruction pointer. The target address may be specified in the instruction or may come from a register value. If the branch branches back to itself a branch target exception will be generated.

Instruction Format: FBcc, BRcc

Cond ₃	Mne.	
0	FBEQ	Rs1 = Rs2
1	FBNE	Rs1 \neq Rs2
2	FBLT / BGT	Rs1 < Rs2 or Rs2 > Rs1
3	FBGE / BLE	Rs1 \geq Rs2 or Rs2 \leq Rs1
4	reserved	
5	reserved	
6	reserved	
7	FBUN	Rs1 ?? Rs2 (branch if unordered)

Clock Cycles: 0.33 (Ideal), 1 Average

Execution Units: FCU Only

Exceptions: branch target

FCLASS – Classify Value

Description:

FCLASS classifies the value in a floating-point register Rs1 and returns the information as a bit vector in the integer register Rd.

Bit	Meaning
0	1 = negative infinity
1	1 = negative number
2	1 = negative subnormal number
3	1 = negative zero
4	1 = positive zero
5	1 = positive subnormal number
6	1 = positive number
7	1 = positive infinity
8	1 = signalling nan
9	1 = quiet nan

FCMP - Float Compare

Description:

The register compare instruction compares two registers or a register and an immediate constant as floating-point values and sets the flags in the target register as a result.

Instruction Format: FLT2, FLTI

Clock Cycles: 2

Execution Units: Floating Point

Operation:

```
if Rs1 < Rs2
    Rd[1] = true
else
    Rd[1] = false
if mag Rs1 < mag Rs2
    Rd[2] = true
else
    Rd[2] = false
if Rs1 = Rs2
    Rd[0] = true
else
    Rd[0] = false
if Rs1 <= Rs2
    Rd[3] = true
else
    Rd[3] = false
if unordered
    Rd[4] = true
else
    Rd[4] = false
```

FCX – Clear Floating-Point Exceptions

Description:

This instruction clears floating point exceptions. The Exceptions to clear are identified as the bits set in the union of register Rs1 and an immediate field in the instruction. Either the immediate or Rs1 should be zero.

Instruction Format: FLT2

Execution Units: All Floating Point

Operation:

Exceptions:

Bit	Exception Enabled
0	global invalid operation clears the following: <ul style="list-style-type: none"> - division of infinities - zero divided by zero - subtraction of infinities - infinity times zero - NaN comparison - division by zero
1	overflow
2	underflow
3	divide by zero
4	inexact operation
5	summary exception

FDX – Floating Disable Exceptions

Description:

This instruction disables floating point exceptions. The Exceptions disabled are identified as the bits set in the union of register Rs1 and an immediate field in the instruction. Either the immediate or Rs1 should be zero. Exceptions won't be disabled until the instruction commits and the state of the machine is updated. This instruction should be followed by a synchronization instruction (FSYNC) to ensure that following floating point operations recognize the disabled exceptions.

Instruction Format: FXX**Clock Cycles:** 2**Execution Units:** Floating Point

FDIV – Floating point divide

Description:

Divide two floating point numbers in registers Rs2 and Rs2 or an immediate constant and place the result into target register Rd.

Instruction Format: FLT2, FLTI

Clock Cycles: 150 (est).

Execution Units: Floating Point

FEX – Floating Enable Exceptions

Description:

This instruction enables floating point exceptions. The Exceptions enabled are identified as the bits set in the union of register Rs1 and an immediate field in the instruction. Either the immediate or Rs1 should be zero. Exceptions won't be enabled until the instruction commits and the state of the machine is updated. This instruction should be followed by a synchronization instruction (FSYNC) to ensure that following floating point operations recognize the enabled exceptions.

Instruction Format: FXX

Clock Cycles: 2

Execution Units: Floating Point

FMA – Floating Point Multiply and Add

Description:

Multiply two floating point numbers in registers Rs1 and Rs2 or an immediate constant then add the floating-point number in Rs3, and place the result into target register Rd.

Instruction Format: FMA

Clock Cycles: 18

Execution Units: Floating Point

FMAX – Maximum Value

Description:

Determines the maximum of three values in registers Rs1, Rs2, Rs3 and places the result in the target register Rd. In order to obtain the maximum value of two registers specify one of the registers twice.

Instruction Format: FLT3

Clock Cycles: 0.33

Execution Units: Integer

Operation:

```
IF Rs1 > Rs2 and Rs1 > Rs3
    Rd = Rs1
else if Rs2 > Rs3
    Rd = Rs2
else
    Rd = Rs3
```

Exceptions: none

FMIN – Minimum Value

Description:

Determines the minimum of three values in registers Rs1, Rs2, Rs3 and places the result in the target register Rd. In order to obtain the minimum value of two registers specify one of the registers twice.

Instruction Format: FLT3

Clock Cycles: 0.33

Execution Units: Integer

Operation:

```
IF Rs1 < Rs2 and Rs1 < Rs3
    Rd = Rs1
else if Rs2 < Rs3
    Rd = Rs2
else
    Rd = Rs3
```

Exceptions: none

FMS – Floating Point Multiply and Subtract

Description:

Multiply two floating point numbers in registers Rs1 and Rs2 or an immediate constant then subtract the floating-point number in Rs3, and place the result into target register Rd.

Instruction Format: FMA

Clock Cycles: 18

Execution Units: Floating Point

FMUL – Floating point multiplication

Description:

Multiply two floating point numbers in registers Rs1 and Rs2 or an immediate constant and place the result into target register Rd.

Instruction Format: FLT2, FLTI

Clock Cycles: 6

Execution Units: Floating Point

FNABS – Floating Negative Absolute Value

Description:

Take the negative absolute value of the floating-point number in register Rs1 and place the result into target register Rd. The sign bit (bit 79) of the register is set to one. No rounding of the number occurs.

Instruction Format: FLT1

Clock Cycles: 0.33

Execution Units: Floating Point

FNEG – Floating Negative Value

Description:

Negate the value of the floating-point number in register Rs1 and place the result into target register Rd. The sign bit (bit 79) of the register is inverted. No rounding of the number occurs.

Instruction Format: FLT1

Clock Cycles: 0.33

Execution Units: Floating Point

FNEXT – Floating Point Next Value

Description:

FNEXT returns the next representable value in the direction indicated by the difference between the values in Rs1 and Rs2.

Instruction Format: FLT2

Clock Cycles: 1

Execution Units: Floating Point

FNMA – Floating Point Negate Multiply and Add

Description:

Multiply two floating point numbers in registers Rs1 and Rs2 or an immediate constant negate the result then add the floating-point number in Rs3, and place the result into target register Rd.

Instruction Format: FMA

Clock Cycles: 18

Execution Units: Floating Point

FNMS – Floating Point Negate Multiply and Subtract

Description:

Multiply two floating point numbers in registers Rs1 and Rs2 or an immediate constant negate the result then subtract the floating-point number in Rs3, and place the result into target register Rd.

Instruction Format: FMA

Clock Cycles: 18

Execution Units: Floating Point

FRES – Reciprocal Estimate

Description:

This function uses a 1024 entry 16-bit precision lookup table to create a piece-wise approximation of the reciprocal and linear interpolation to approximate the reciprocal of the value in Rs1. The value is returned in Rd as an 80-bit floating-point value. The value returned is accurate to about eight bits.

Instruction Format: FLT1A

Clock Cycles: 5

Execution Units: Floating Point

FRSQRTE – Float Reciprocal Square Root Estimate

Description:

Estimate the reciprocal of the square root of the number in register Rs1 and place the result into target register Rd. This instruction may not be present in all models of the core.

Instruction Format: FLT1A

Clock Cycles: 3

Execution Units: Floating Point

Notes:

The estimate is only accurate to about 3%. The estimate is performed in single precision (32-bit) floating point, then converted to an 80-bit format. That means that input values must be in the range of a 32-bit floating point number. Values outside of this range will return infinity or zero as a result.

Taking the reciprocal square root of a negative number results in a Nan output.

FSIGN – Floating Sign

Description:

FSIGN returns a value indicating the sign of the floating-point number. If the value is zero, the target register is set to zero. If the value is negative the target register is set to the floating-point value -1.0. Otherwise the target register is set to the floating-point value +1.0. No rounding of the result occurs.

Instruction Format: FLT1

Clock Cycles: 0.33

Execution Units: Floating Point

FSEQ - Float Set if Equal

Description:

The register compare instruction compares two registers or a register and an immediate constant as floating-point values for equality and sets the integer target register as a result. Note that negative and positive zero are considered equal.

Instruction Format: FLT2, FLTI

Clock Cycles: 0.33

Execution Units: Floating Point

Operation:

```
if Rs1 == Rs2
    Rd = true
else
    Rd = false
```

FSGE - Float Set if Greater Than or Equal

Description:

The register compare instruction compares two registers or a register and an immediate constant as floating-point values for greater than or equal and sets the integer target register as a result.

Instruction Format: FLT2, FLTI

Clock Cycles: 0.33

Execution Units: Floating Point

Operation:

```
if Rs1 >= Rs2
    Rd = true
else
    Rd = false
```


FSGN[op] – Sign Operation

Description:

FSGNCPY copies the sign of one value to another value. This instruction may be used to get the absolute value.

FSGNINV copies the inverted sign of one value to another. This instruction may be used to get the negative absolute value.

FSGNOR logically or's the sign of the two values.

FSGNAND logically and's the sign of the two values.

FSGNXOR logically xor's the sign of the two values.

FSGNXNOR logically xnor's the sign of the two values.

Instruction Format: FLT2, FLTI

Op ₃	Operation
0	FSGNCPY
1	FSGNINV
4	FSGNAND
5	FSGNOR
6	FSGNXOR
7	FSGNXNOR

Clock Cycles: 0.33

Execution Units: Floating Point

Operation:

FSGT - Float Set if Greater Than

Description:

The register compare instruction compares two registers or a register and an immediate constant as floating-point values for greater than and sets the integer target register as a result.

Instruction Format: FLT2, FLTI

Clock Cycles: 0.33

Execution Units: Floating Point

Operation:

```
if Rs1 > Rs2
    Rd = true
else
    Rd = false
```

FSIG – Sigmoid Approximate

Description:

This function uses a 1024 entry 32-bit precision lookup table with linear interpolation to approximate the logistic sigmoid function in the range -8.0 to +8.0. Outside of this range 0.0 or +1.0 is returned. The sigmoid output is between 0.0 and +1.0. The value of the sigmoid for register Rs1 is returned in register Rd as an 80-bit floating-point value.

Instruction Format: FLT1A

Clock Cycles: 5

Execution Units: Floating Point

FSLE - Float Set if Less Than or Equal

Description:

The register compare instruction compares two registers or a register and an immediate constant as floating-point values for less than or equal and sets the integer target register as a result.

Instruction Format: FLT2, FLTI

Clock Cycles: 0.33

Execution Units: Floating Point

Operation:

```
if Rs1 <= Rs2
    Rd = true
else
    Rd = false
```

FSLT - Float Set if Less Than

Description:

The register compare instruction compares two registers or a register and an immediate constant as floating-point values for less than and sets the integer target register as a result.

Instruction Format: FLT2, FLTI

Clock Cycles: 0.33

Execution Units: Floating Point

Operation:

```
if Rs1 < Rs2
    Rd = true
else
    Rd = false
```

FSNE - Float Set if Not Equal

Description:

The register compare instruction compares two registers or a register and an immediate constant as floating-point values for inequality and sets the integer target register as a result. Note that negative and positive zero are considered equal.

Instruction Format: FLT2, FLTI

Clock Cycles: 0.33

Execution Units: Floating Point

Operation:

```
if Rs1 <> Rs2
    Rd= true
else
    Rd= false
```

FSQRT – Floating point square root

Description:

Take the square root of the floating-point number in register Rs1 and place the result into target register Rd. The sign bit (bit 79) of the register is set to zero. This instruction can generate NaNs.

Instruction Format: FLT1

Clock Cycles: 150 (est).

Execution Units: Floating Point

FSUB – Floating point subtraction

Description:

Subtract two floating-point numbers in registers Rs1 and Rs2 or an immediate constant and place the result into target register Rd.

Instruction Format: FLT2, FLTI

Clock Cycles: 10

Execution Units: Floating Point

FSUN - Float Set if Unordered

Description:

The register compare instruction compares two registers as floating-point values for unorderedness and sets the target register as a result. The source operands are from the register set specified in the FSTAT CSR. The target register is from the register set specified in the MSTATUS CSR.

Instruction Format: FLT2

Clock Cycles: 0.33

Execution Units: Floating Point

Operation:

```
if Ra ?? Rb
    Rt = true
else
    Rt = false
```

FSYNC -Synchronize

Description:

All floating-point instructions before the FSYNC are completed and committed to the architectural state before floating-point instructions after the FSYNC are issued. This instruction is used to ensure that the machine state is valid before subsequent instructions are executed.

Instruction Format: FSYNC**Clock Cycles:** varies depending on queue contents

FTOI – Floating Convert to Integer

Description:

Convert the floating-point value in Rs1 into an integer and place the result into target register Rd. If the result overflows the value placed in Rd is a maximum integer value.

Instruction Format: FLT1**Clock Cycles:** 3**Execution Units:** Floating Point

FTRUNC – Truncate Value

Description:

The FTRUNC instruction truncates off the fractional portion of the number leaving only a whole value. For instance, trunc(1.5) equals 1.0. Ftrunc does not change the representation of the number. To convert a value to an integer in a fixed-point representation see the FTOI instruction.

Instruction Format: F1

Clock Cycles: 0.33

Execution Units: Floating Point

FTX – Trigger Floating Point Exceptions

Description:

This instruction triggers floating point exceptions. The Exceptions to trigger are identified as the bits set in the union of register Rs1 and an immediate field in the instruction. Either the immediate or Rs1 should be zero.

Instruction Format: FXX

Execution Units: Floating Point

Operation:

Exceptions:

Bit	Exception Enabled
0	global invalid operation
1	overflow
2	underflow
3	divide by zero
4	inexact operation
5	reserved

ISNAN – Is Not a Number

Description:

Test the value in Rs1 to see if it's a nan (not a number) and return true (1) or false (0) in a general-purpose register Rd.

Instruction Format: FLT1

Clock Cycles: 0.33

Execution Units: Floating Point

Example:

```
isnan  $r1,$fp7
```

ITOF – Convert Integer to Float

Description:

Convert the integer value in Rs1 into a floating-point value and place the result into target register Rd. Rs1 is from either the floating-point register set or the integer register set, Rd is in the floating-point register set. Some precision of the integer converted may be lost if the integer is larger than 64 bits. Extended double precision floating point values only have a precision of 65 bits.

Instruction Format: FLT1

M1: Ra register set. 0 = use float register set, 1 = move from integer register set to float

Clock Cycles: 3

Execution Units: Floating Point

Glossary

Burst Access

A burst access is a number of bus accesses that occur rapidly in a row in a known sequence. If hardware supports burst access the cycle time for access to the device is drastically reduced. For instance dynamic RAM memory access is really fast for sequential burst access, and somewhat slower for random access.

BTB

An acronym for Branch Target Buffer. The branch target buffer is used to improve the performance of a processing core. The BTB is a table that stores the branch target from previously executed branch instructions. A typical table may contain 1024 entries. The table is typically indexed by part of the branch address. Since the target address of a branch type instruction may not be known at fetch time, the address is speculated to be the address in the branch target buffer. This allows the machine to fetch instructions in a continuous fashion without pipeline bubbles. In many cases the calculated branch address from a previously executed instruction remains the same the next time the same instruction is executed. If the address from the BTB turns out to be incorrect, then the machine will have to flush the instruction queue or pipeline and begin fetching instructions from the correct address.

FPGA

An acronym for Field Programmable Gate Array. FPGA's consist of a large number of small RAM tables, flip-flops and other logic. These are all connected together with a programmable connection network. FPGA's are 'in the field' programmable, and usually re-programmable. An FPGA's re-programmability is typically RAM based. They are often used with configuration PROM's so they may be loaded to perform specific functions.

HDL

An acronym that stands for 'Hardware Description Language'. A hardware description language is used to describe hardware constructs at a high level.

Instruction Bundle

A group of instructions. It is sometimes required to group instructions together into bundle. For instance, all instructions in a bundle may be executed simultaneously on a processor as a unit. Instructions may also need to be grouped if they are oddball in size for example 41 bits, so that they can be fit evenly into memory. Typically, a bundle has some bits that are global to the bundle, such as template bits, in addition to the encoded instructions.

Instruction Pointers

A processor register dedicated to addressing instructions in memory. It is also often called a program counter. The program counter got its name because it usually increments (or counts) automatically after an instruction is fetched. In early machines in some rare cases the program counter did not count in a sequential binary fashion, but

instead used other forms of a counter such as a grey counter or linear feedback shift register. In some machines the program counter addresses bundles of instructions rather than individual instructions. This is common with some stack machines where multiple instructions are packed into a memory word.

ISA

An acronym for Instruction Set Architecture. The group of instructions that an architecture supports. ISA's are sometimes categorized at extreme edges as RISC or CISC. FT64 falls somewhere in between with features of both RISC and CISC architectures.

Linear Address

A linear address is the resulting address from a virtual address after segmentation has been applied.

Physical Address

A physical address is the final address seen by the memory system after both segmentation and paging have been applied to a virtual address. One can think of a physical address as one that is "physically" wired to the memory.

Program Counter

A processor register dedicated to addressing instructions in memory. It is also often and perhaps more aptly called an instruction pointer. The program counter got its name because it usually increments (or counts) automatically after an instruction is fetched. In early machines in some rare cases the program counter did not count in a sequential binary fashion, but instead used other forms of a counter such as a grey counter or linear feedback shift register. In some machines the program counter addresses bundles of instructions rather than individual instructions. This is common with some stack machines where multiple instructions are packed into a memory word.

ROB

An acronym for ReOrder Buffer. The re-order buffer allows instructions to execute out of order yet update the machine's state in order by tracking instruction state and variables. In FT64 the re-order buffer is a circular queue with a head and tail pointers. Instructions at the head are committed if done to the machine's state then the head advanced. New instructions are queued at the buffer's tail as long as there is room in the queue. Instructions in the queue may be processed out of the order that they entered the queue in depending on the availability of resources (register values and functional units).

RSB

An acronym that stands for return stack buffer. A buffer of addresses used to predict the return address which increases processor performance. The RSB is usually small, typically 16 entries. When a return instruction is detected at time of fetch the RSB is accessed to determine the address of the next instruction to fetch. Predicting the return address allows the processing core to continuously fetch instructions in a speculative fashion without bubbles in the pipeline. The return address in the RSB may turn out to be detected as incorrect during execution of the return instruction, in which case the pipeline

or instruction queue will need to be flushed and instructions fetched from the proper address.

SIMD

An acronym that stands for ‘Single Instruction Multiple Data’. SIMD instructions are usually implemented with extra wide registers. The registers contain multiple data items, such as a 128-bit register containing four 32-bit numbers. The same instruction is applied to all the data items in the register at the same time. For some applications SIMD instructions can enhance performance considerably.

Stack Pointer

A processor register dedicated to addressing stack memory. Sometimes this register is assigned by convention from the general register pool. This register may also sometimes index into a small dedicated stack memory that is not part of the main memory system. Sometimes machines have multiple stack pointers for different purposes, but they all work on the idea of a stack. For instance, in Forth machines there are typically two stacks, one for data and one for return addresses.

TLB

TLB stands for translation look-aside buffer. This buffer is used to store address translations for fast memory access in a system with an mmu capable of performing address translations.

Miscellaneous

Reference Material

Below is a short list of some of the reading material the author has studied. The author has downloaded a fair number of documents on computer architecture from the web. Too many to list.

Modern Processor Design Fundamentals of Superscalar Processors by John Paul Shen, Mikko H. Lipasti. Waveland Press, Inc.

Computer Architecture A Quantitative Approach, Second Edition, by John L Hennessy & David Patterson, published by Morgan Kaufman Publishers, Inc. San Francisco, California is a good book on computer architecture. There is a newer edition of the book available.

Memory Systems Cache, DRAM, Disk by Bruce Jacob, Spencer W. Ng., David T. Wang, Samuel Rodriguez, Morgan Kaufman Publishers

PowerPC Microprocessor Developer's Guide, SAMS publishing. 201 West 103rd Street, Indianapolis, Indiana, 46290

80386/80486 Programming Guide by Ross P. Nelson, Microsoft Press

Programming the 286, C. Vieillefond, SYBEX, 2021 Challenger Drive #100, Alameda, CA 94501

Tech. Report UMD-SCA-2000-02 ENEE 446: Digital Computer Design — An Out-of-Order RiSC-16

Programming the 65C816, David Eyes and Ron Lichty, Western Design Centre Inc.

Microprocessor Manuals from Motorola, and Intel,

The SPARC Architecture Manual Version 8, SPARC International Inc, 535 Middlefield Road. Suite 210 Menlo Park California, CA 94025

The SPARC Architecture Manual Version 9, SPARC International Inc, San Jose California, PTR Prentice Hall, Englewood Cliffs, New Jersey, 07632

The MMIX processor: <http://mmix.cs.hm.edu/doc/instructions-en.html>

RISCV 2.0 Spec, Andrew Waterman, Yunsup Lee, David Patterson, Krste Asanović CS Division, EECS Department, University of California, Berkeley {waterman|yunsup|pattsrn|krste}@eecs.berkeley.edu

Trademarks

IBM® is a registered trademark of International Business Machines Corporation. Intel® is a registered trademark of Intel Corporation. HP® is a registered trademark of Hewlett-Packard Development Company. "SPARC® is a registered trademark of SPARC International, Inc.

WISHBONE Compatibility Datasheet

The NVIO core may be directly interfaced to a WISHBONE compatible bus.

WISHBONE Datasheet		
WISHBONE SoC Architecture Specification, Revision B.3		
Description:	Specifications:	
General Description:	Central processing unit (CPU core)	
Supported Cycles:	MASTER, READ / WRITE MASTER, READ-MODIFY-WRITE MASTER, BLOCK READ / WRITE, BURST READ (FIXED ADDRESS)	
Data port, size:	128 bit	
Data port, granularity:	8 bit	
Data port, maximum operand size:	128 bit	
Data transfer ordering:	Little Endian	
Data transfer sequencing	any (undefined)	
Clock frequency constraints:		
Supported signal list and cross reference to equivalent WISHBONE signals	Signal Name:	WISHBONE Equiv.
	ack_i	ACK_I
	adr_o(31:0)	ADR_O()
	clk_i	CLK_I
	dat_i(127:0)	DAT_I()
	dat_o(127:0)	DAT_O()
	cyc_o	CYC_O
	stb_o	STB_O
	wr_o	WE_O
	sel_o(15:0)	SEL_O
	cti_o(2:0)	CTI_O
	bte_o(1:0)	BTE_O
Special Requirements:		